

1. Business Problem

1.1 Problem Description

Netflix is all about connecting people to the movies they love. To help customers find those movies, they developed world-class movie recommendation system: CinematchSM. Its job is to predict whether someone will enjoy a movie based on how much they liked or disliked other movies. Netflix use those predictions to make personal movie recommendations based on each customer's unique tastes. And while **Cinematch** is doing pretty well, it can always be made better.

Now there are a lot of interesting alternative approaches to how Cinematch works that netflix haven't tried. Some are described in the literature, some aren't. We're curious whether any of these can beat Cinematch by making better predictions. Because, frankly, if there is a much better approach it could make a big difference to our customers and our business.

Credits: <https://www.netflixprize.com/rules.html>

1.2 Problem Statement

Netflix provided a lot of anonymous rating data, and a prediction accuracy bar that is 10% better than what Cinematch can do on the same training data set. (Accuracy is a measurement of how closely predicted ratings of movies match subsequent actual ratings.)

1.3 Sources

- <https://www.netflixprize.com/rules.html>
- <https://www.kaggle.com/netflix-inc/netflix-prize-data>
- Netflix blog: <https://medium.com/netflix-techblog/netflix-recommendations-beyond-the-5-stars-part-1-55838468f429> (very nice blog)
- surprise library: <http://surpriselib.com/> (we use many models from this library)
- surprise library doc: http://surprise.readthedocs.io/en/stable/getting_started.html (we use many models from this library)
- installing surprise: <https://github.com/NicolasHug/ Surprise#installation>
- Research paper: <http://courses.ischool.berkeley.edu/i290-dm/s11/SECURE/a1-koren.pdf> (most of our work was inspired by this paper)
- SVD Decomposition : <https://www.youtube.com/watch?v=P5mlg91as1c>

1.4 Real world/Business Objectives and constraints

Objectives:

1. Predict the rating that a user would give to a movie that he has not yet rated.
2. Minimize the difference between predicted and actual rating (RMSE and MAPE)

Constraints:

1. Some form of interpretability.

2. Machine Learning Problem

2.1 Data

2.1.1 Data Overview

Get the data from : <https://www.kaggle.com/netflix-inc/netflix-prize-data/data>

Data files :

- combined_data_1.txt
- combined_data_2.txt
- combined_data_3.txt
- combined_data_4.txt
- movie_titles.csv

The first line of each file [combined_data_1.txt, combined_data_2.txt, combined_data_3.txt, combined_data_4.txt] contains the movie id followed by a colon. Each subsequent line in the file corresponds to a rating from a customer and its date in the following format:

CustomerID,Rating,Date

MovieIDs range from 1 to 17770 sequentially.

CustomerIDs range from 1 to 2649429, with gaps. There are 480189 users.

Ratings are on a five star (integral) scale from 1 to 5.

Dates have the format YYYY-MM-DD.

2.1.2 Example Data point

1:

1488844,3,2005-09-06
822109,5,2005-05-13
885013,4,2005-10-19
30878,4,2005-12-26
823519,3,2004-05-03
893988,3,2005-11-17
124105,4,2004-08-05
1248029,3,2004-04-22
1842128,4,2004-05-09
2238063,3,2005-05-11
1503895,4,2005-05-19
2207774,5,2005-06-06
2590061,3,2004-08-12
2442,3,2004-04-14
543865,4,2004-05-28
1209119,4,2004-03-23
804919,4,2004-06-10
1086807,3,2004-12-28
1711859,4,2005-05-08
372233,5,2005-11-23
1080361,3,2005-03-28
1245640,3,2005-12-19
558634,4,2004-12-14
2165002,4,2004-04-06
1181550,3,2004-02-01
1227322,4,2004-02-06
427928,4,2004-02-26
814701,5,2005-09-29
808731,4,2005-10-31
662870,5,2005-08-24
337541,5,2005-03-23
786312,3,2004-11-16
1133214,4,2004-03-07
1537427,4,2004-03-29
1209954,5,2005-05-09
2381599,3,2005-09-12
525356,2,2004-07-11
1910569,4,2004-04-12
2263586,4,2004-08-20
2421815,2,2004-02-26
1009622,1,2005-01-19
1481961,2,2005-05-24
401047,4,2005-06-03
2179073,3,2004-08-29
1434636,3,2004-05-01
93986,5,2005-10-06
1308744,5,2005-10-29
2647871,4,2005-12-30
1905581,5,2005-08-16
2508819,3,2004-05-18
1578279,1,2005-05-19
1159695,4,2005-02-15
2588432,3,2005-03-31

2.2 Mapping the real world problem to a Machine Learning Problem

2.2.1 Type of Machine Learning Problem

For a given movie and user we need to predict the rating would be given by him/her to the movie.

The given problem is a Recommendation problem

It can also be seen as a Regression problem

2.2.2 Performance metric

- Mean Absolute Percentage Error: https://en.wikipedia.org/wiki/Mean_absolute_percentage_error
- Root Mean Square Error: https://en.wikipedia.org/wiki/Root-mean-square_deviation

2.2.3 Machine Learning Objective and Constraints

1. Minimize RMSE.
2. Try to provide some interpretability.

```
In [1]: # this is just to know how much time will it take to run this entire ipython notebook
ok
from datetime import datetime
# globalstart = datetime.now()
import pandas as pd
import numpy as np
import matplotlib
matplotlib.use('nbagg')

import matplotlib.pyplot as plt
plt.rcParams.update({'figure.max_open_warning': 0})

import seaborn as sns
sns.set_style('whitegrid')
import os
from scipy import sparse
from scipy.sparse import csr_matrix

from sklearn.decomposition import TruncatedSVD
from sklearn.metrics.pairwise import cosine_similarity
import random
import scipy.stats as stats
plt.style.use('seaborn')
#print(plt.style.available)
from IPython.display import clear_output
```

```
/usr/local/lib/python2.7/dist-packages/ipykernel_launcher.py:6: UserWarning:
This call to matplotlib.use() has no effect because the backend has already
been chosen; matplotlib.use() must be called *before* pylab, matplotlib.pyplot,
or matplotlib.backends is imported for the first time.
```

The backend was *originally* set to 'module://ipykernel.pylab.backend_inline' by the following code:

```
File "/usr/lib/python2.7/runpy.py", line 174, in _run_module_as_main
    "__main__", fname, loader, pkg_name)
File "/usr/lib/python2.7/runpy.py", line 72, in _run_code
    exec code in run_globals
File "/usr/local/lib/python2.7/dist-packages/ipykernel_launcher.py", line 16,
in <module>
    app.launch_new_instance()
File "/usr/local/lib/python2.7/dist-packages/traitlets/config/application.py",
line 657, in launch_instance
    app.initialize(argv)
File "</usr/local/lib/python2.7/dist-packages/decorator.pyc:decorator-gen-121>
", line 2, in initialize
    File "/usr/local/lib/python2.7/dist-packages/traitlets/config/application.py",
line 87, in catch_config_error
    return method(app, *args, **kwargs)
File "/usr/local/lib/python2.7/dist-packages/ipykernel/kernelapp.py", line 462
, in initialize
    self.init_gui_pylab()
File "/usr/local/lib/python2.7/dist-packages/ipykernel/kernelapp.py", line 403
, in init_gui_pylab
    InteractiveShellApp.init_gui_pylab(self)
File "/usr/local/lib/python2.7/dist-packages/IPython/core/shellapp.py", line 2
13, in init_gui_pylab
    r = enable(key)
File "/usr/local/lib/python2.7/dist-packages/IPython/core/interactiveshell.py"
, line 2950, in enable_matplotlib
    pt.activate_matplotlib(backend)
File "/usr/local/lib/python2.7/dist-packages/IPython/core/pylabtools.py", line
309, in activate_matplotlib
    matplotlib.pyplot.switch_backend(backend)
File "/usr/local/lib/python2.7/dist-packages/matplotlib/pyplot.py", line 231,
in switch_backend
    matplotlib.use(newbackend, warn=False, force=True)
File "/usr/local/lib/python2.7/dist-packages/matplotlib/__init__.py", line 142
2, in use
    reload(sys.modules['matplotlib.backends'])
File "/usr/local/lib/python2.7/dist-packages/matplotlib/backends/__init__.py",
line 17, in <module>
    line for line in traceback.format_stack()
```

3. Exploratory Data Analysis

3.1 Preprocessing

3.1.1 Converting / Merging whole data to required format: u_i, m_j, r_ij

```
In [13]: start = datetime.now()
if not os.path.isfile('gdrive/My Drive/Colab Notebooks/data.csv'):
    # Create a file 'data.csv' before reading it
    # Read all the files in netflix and store them in one big file('data.csv')
    # We re reading from each of the four files and appendig each rating to a globa
l file 'train.csv'
    data = open('data.csv', mode='w')

    row = list()
    files=['data_folder/combined_data_1.txt','data_folder/combined_data_2.txt',
           'data_folder/combined_data_3.txt', 'data_folder/combined_data_4.txt']
    for file in files:
        print("Reading ratings from {}".format(file))
        with open(file) as f:
            for line in f:
                del row[:] # you don't have to do this.
                line = line.strip()
                if line.endswith(':'):
                    # All below are ratings for this movie, until another movie app
ears.

                    movie_id = line.replace(':', '')
                else:
                    row = [x for x in line.split(',')]
                    row.insert(0, movie_id)
                    data.write(','.join(row))
                    data.write('\n')
            print("Done.\n")
        data.close()
print('Time taken :', datetime.now() - start)
```

Time taken : 0:00:00.001944

```
In [14]: print("creating the dataframe from data.csv file..")
df = pd.read_csv('gdrive/My Drive/Colab Notebooks/data.csv', sep=',',
                 names=['movie', 'user', 'rating', 'date'])
df.date = pd.to_datetime(df.date)
print('Done.\n')

# we are arranging the ratings according to time.
print('Sorting the dataframe by date..')
df.sort_values(by='date', inplace=True)
print('Done..')
```

creating the dataframe from data.csv file..
Done.

Sorting the dataframe by date..
Done..


```
In [15]: df.head()
```

```
Out[15]:
```

	movie	user	rating	date
56431994	10341	510180	4	1999-11-11
9056171	1798	510180	5	1999-11-11
58698779	10774	510180	3	1999-11-11
48101611	8651	510180	2	1999-11-11
81893208	14660	510180	2	1999-11-11

```
In [16]: df.describe()['rating']
```

```
Out[16]: count    1.004805e+08  
mean        3.604290e+00  
std         1.085219e+00  
min         1.000000e+00  
25%         3.000000e+00  
50%         4.000000e+00  
75%         4.000000e+00  
max         5.000000e+00  
Name: rating, dtype: float64
```

3.1.2 Checking for NaN values

```
In [0]: # just to make sure that all Nan containing rows are deleted..  
print("No of Nan values in our dataframe : ", sum(df.isnull().any()))
```

```
No of Nan values in our dataframe : 0
```

3.1.3 Removing Duplicates

```
In [0]: dup_bool = df.duplicated(['movie','user','rating'])  
dups = sum(dup_bool) # by considering all columns..( including timestamp)  
print("There are {} duplicate rating entries in the data..".format(dups))
```

```
There are 0 duplicate rating entries in the data..
```

3.1.4 Basic Statistics (#Ratings, #Users, and #Movies)

```
In [0]: print("Total data ")
print("-"*50)
print("\nTotal no of ratings :",df.shape[0])
print("Total No of Users      :", len(np.unique(df.user)))
print("Total No of movies     :", len(np.unique(df.movie)))
```

Total data

```
Total no of ratings : 100480507
Total No of Users   : 480189
Total No of movies  : 17770
```

3.2 Splitting data into Train and Test(80:20)

```
In [0]: if not os.path.isfile('train.csv'):
        # create the dataframe and store it in the disk for offline purposes..
        df.iloc[:int(df.shape[0]*0.80)].to_csv("train.csv", index=False)

        if not os.path.isfile('test.csv'):
            # create the dataframe and store it in the disk for offline purposes..
            df.iloc[int(df.shape[0]*0.80):].to_csv("test.csv", index=False)

train_df = pd.read_csv("train.csv", parse_dates=['date'])
test_df = pd.read_csv("test.csv")
```

3.2.1 Basic Statistics in Train data (#Ratings, #Users, and #Movies)

```
In [4]: # movies = train_df.movie.value_counts()
# users = train_df.user.value_counts()
print("Training data ")
print("-"*50)
print("\nTotal no of ratings :",train_df.shape[0])
print("Total No of Users      :", len(np.unique(train_df.user)))
print("Total No of movies     :", len(np.unique(train_df.movie)))
```

Training data

```
('Total no of ratings :', 80384405)
('Total No of Users   :', 405041)
('Total No of movies  :', 17424)
```

3.2.2 Basic Statistics in Test data (#Ratings, #Users, and #Movies)

```
In [5]: print("Test data ")
print("-"*50)
print("\nTotal no of ratings :",test_df.shape[0])
print("Total No of Users      :", len(np.unique(test_df.user)))
print("Total No of movies     :", len(np.unique(test_df.movie)))
```

Test data

```
('Total no of ratings :', 20096102)
('Total No of Users   :', 349312)
('Total No of movies  :', 17757)
```

3.3 Exploratory Data Analysis on Train data

```
In [0]: # method to make y-axis more readable
def human(num, units = 'M'):
    units = units.lower()
    num = float(num)
    if units == 'k':
        return str(num/10**3) + " K"
    elif units == 'm':
        return str(num/10**6) + " M"
    elif units == 'b':
        return str(num/10**9) + " B"
```

3.3.1 Distribution of ratings

```
In [7]: fig, ax = plt.subplots()
plt.title('Distribution of ratings over Training dataset', fontsize=15)
sns.countplot(train_df.rating)
ax.set_yticklabels([human(item, 'M') for item in ax.get_yticks()])
ax.set_ylabel('No. of Ratings(Millions)')

plt.show()
```

/usr/local/lib/python2.7/dist-packages/seaborn/categorical.py:1428: FutureWarning: remove_na is deprecated and is a private function. Do not use.
stat_data = remove_na(group_data)



Add new column (week day) to the data set for analysis.

```
In [10]: # It is used to skip the warning 'SettingWithCopyWarning'..
pd.options.mode.chained_assignment = None # default='warn'

train_df['day_of_week'] = train_df.date.dt.weekday_name

train_df.tail()
```

Out[10]:

	movie	user	rating	date	day_of_week
80384400	12074	2033618	4	2005-08-08	Monday
80384401	862	1797061	3	2005-08-08	Monday
80384402	10986	1498715	5	2005-08-08	Monday
80384403	14861	500016	4	2005-08-08	Monday
80384404	5926	1044015	5	2005-08-08	Monday

```
In [0]: data_by_day_m = train_df.groupby('day_of_week').mean()
data_by_day_c = train_df.groupby('day_of_week').count()
```

```
In [12]: data_by_day_m.loc['Monday', 'rating']
```

Out[12]: 3.577249889462102

3.3.1.1 Effect on the rating on day of week

```
In [0]: labels = ('Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday')
data = []
sig = []

for elem in labels:
    a = list(train_df[train_df['day_of_week'] == elem]['rating'].values)
    t, p = stats.ttest_1samp(a, np.mean(train_df['rating'].values))
    if p < 0.001:
        sig.append('Yes')
    else:
        sig.append('No')

    data.append(a)


means = map(np.mean, data)
se = map(stats.sem, data)
n_s = map(len, data)
```

```
In [15]: x_pos = list(range(1,2*len(data),2))
color=['#FF6666', '#FFB266', '#FFFF66', '#B2FF66', '#66B2FF', '#6666FF', '#B266FF']

fig, ax = plt.subplots(figsize = (7.5, 7.5))

plt.plot([0,14],[np.mean(train_df['rating'].values)]*2, linewidth = 1, color = 'k',
linestyle = '--')
plt.bar(x_pos,
        # using the data from the mean_values
        means,
        # with a y-error lines set as standar errors
        yerr=se,
        # aligned in the center
        align='center',
        # with color
        color = color,
        # transparency
        alpha=.75)

ax.set_xticks(x_pos)
ax.set_xticklabels(labels)
plt.ylim(2,5)
plt.xticks(rotation=45)
plt.xlabel('Day of the week')
plt.ylabel('Average rating')
plt.show()
```



```
In [16]: for i in range(len(labels)):
          print(labels[i] + ' average rating : ' + str(np.mean(data[i])))
          print('Was ' + labels[i] + ' different than the overall mean? : ' + sig[i])
          print(' ')
```

Monday average rating : 3.577249889462102
Was Monday different than the overall mean? : Yes

Tuesday average rating : 3.574438350090171
Was Tuesday different than the overall mean? : Yes

Wednesday average rating : 3.583750996692474
Was Wednesday different than the overall mean? : No

Thursday average rating : 3.5824631807848175
Was Thursday different than the overall mean? : No

Friday average rating : 3.585274296026911
Was Friday different than the overall mean? : Yes

Saturday average rating : 3.5917909557481895
Was Saturday different than the overall mean? : Yes

Sunday average rating : 3.5941436926181596
Was Sunday different than the overall mean? : Yes


3.3.1.2 Average user rating

```
In [0]: #this part groups the data by user id applying the mean() function
data_by_user_m = train_df.groupby('user').mean()
#this part groups the data by user id applying the count() function
data_by_user_c = train_df.groupby('user').count()
```

```
In [4]: values, bins, _ = plt.hist(data_by_user_m['rating'].values, bins = 50, normed = True)
        area = sum(np.diff(bins)*values)
        plt.xlabel('Average user rating (Stars)')
        plt.ylabel('Probability Density Function')
        plt.show()
```

/usr/local/lib/python2.7/dist-packages/matplotlib/axes/_axes.py:6571: UserWarning: The 'normed' kwarg is deprecated, and has been replaced by the 'density' kwarg.

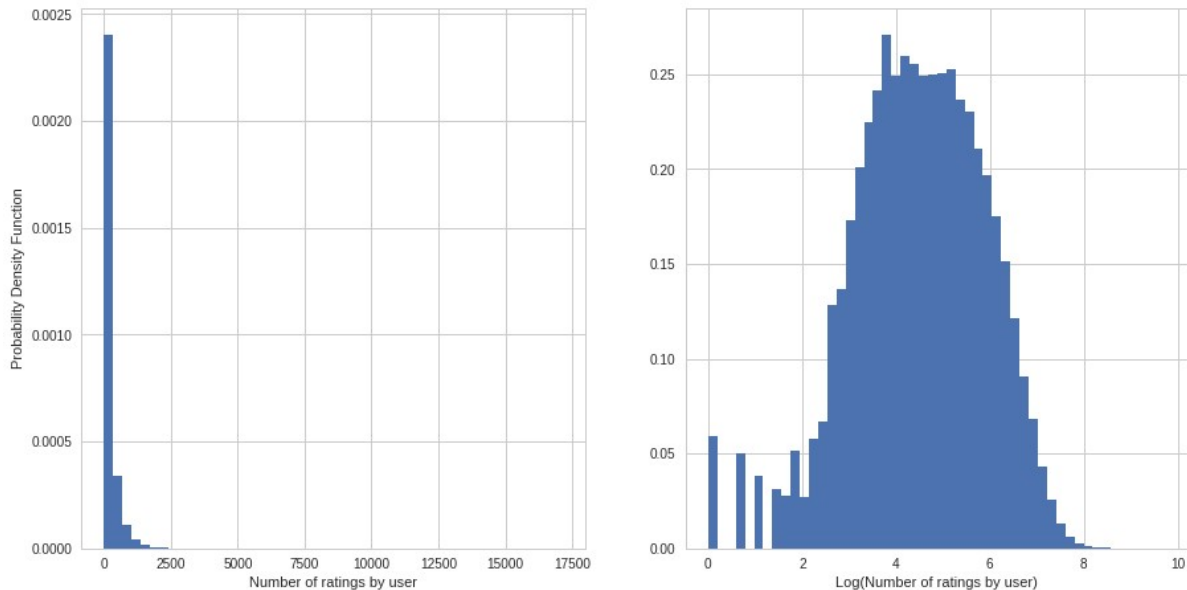
warnings.warn("The 'normed' kwarg is deprecated, and has been "



```
In [5]: fig = plt.figure(figsize=(15,7.5))
plt.subplot(1,2,1)
plt.hist(data_by_user_c['rating'].values, bins = 50, normed = True)
plt.ylabel('Probability Density Function')
plt.xlabel('Number of ratings by user')


plt.subplot(1,2,2)
plt.hist(np.log(data_by_user_c['rating'].values), bins = 50, normed = True)
plt.xlabel('Log(Number of ratings by user)')

plt.show()
```



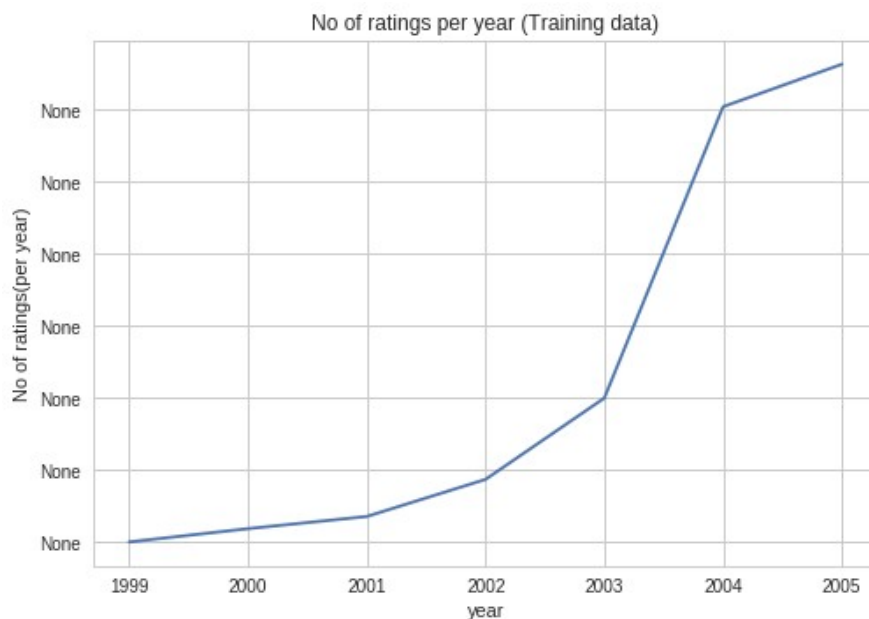
3.3.2 Number of Ratings per a month


```
In [8]: ax = train_df.resample('m', on='date')['rating'].count().plot()
ax.set_title('No of ratings per month (Training data)')
plt.xlabel('Month')
plt.ylabel('No of ratings(per month)')
ax.set_yticklabels([human(item, 'M') for item in ax.get_yticks()])
plt.show()
```



3.3.2.1 Number of Ratings per a year

```
In [9]: ax = train_df.resample('y', on='date')['rating'].count().plot()
ax.set_title('No of ratings per year (Training data)')
plt.xlabel('year')
plt.ylabel('No of ratings(per year)')
ax.set_yticklabels([human(item, 'y') for item in ax.get_yticks()])
plt.show()
```



3.3.3 Analysis on the Ratings given by user

```
In [0]: no_of Rated_movies_per_user = train_df.groupby(by='user')['rating'].count().sort_values(ascending=False)

no_of Rated_movies_per_user.head()
```


```
Out[0]: user
305344      17112
2439493     15896
387418      15402
1639792      9767
1461435      9447
Name: rating, dtype: int64
```

```
In [0]: fig = plt.figure(figsize=plt.figaspect(.5))

ax1 = plt.subplot(121)
sns.kdeplot(no_of Rated_movies_per_user, shade=True, ax=ax1)
plt.xlabel('No of ratings by user')
plt.title("PDF")

ax2 = plt.subplot(122)
sns.kdeplot(no_of Rated_movies_per_user, shade=True, cumulative=True, ax=ax2)
plt.xlabel('No of ratings by user')
plt.title('CDF')

plt.show()
```



```
In [0]: no_of Rated_movies_per_user.describe()
```

```
Out[0]: count      405041.000000  
mean         198.459921  
std          290.793238  
min           1.000000  
25%          34.000000  
50%          89.000000  
75%         245.000000  
max         17112.000000  
Name: rating, dtype: float64
```

There, is something interesting going on with the quantiles..

```
In [0]: quantiles = no_of Rated_movies_per_user.quantile(np.arange(0,1.01,0.01), interpolation='higher')
```


```

In [0]: plt.title("Quantiles and their Values")
quantiles.plot()
# quantiles with 0.05 difference
plt.scatter(x=quantiles.index[::5], y=quantiles.values[::5], c='orange', label="qua
ntiles with 0.05 intervals")
# quantiles with 0.25 difference
plt.scatter(x=quantiles.index[::25], y=quantiles.values[::25], c='m', label = "quan
tiles with 0.25 intervals")
plt.ylabel('No of ratings by user')
plt.xlabel('Value at the quantile')
plt.legend(loc='best')

# annotate the 25th, 50th, 75th and 100th percentile values....
for x,y in zip(quantiles.index[::25], quantiles[::25]):
    plt.annotate(s="({} , {})".format(x,y), xy=(x,y), xytext=(x-0.05, y+500)
                , fontweight='bold')

plt.show()

```



```
In [0]: quantiles[::5]
```

```
Out[0]: 0.00      1
        0.05      7
        0.10     15
        0.15     21
        0.20     27
        0.25     34
        0.30     41
        0.35     50
        0.40     60
        0.45     73
        0.50     89
        0.55    109
        0.60    133
        0.65    163
        0.70    199
        0.75    245
        0.80    307
        0.85    392
        0.90    520
        0.95    749
        1.00   17112
        Name: rating, dtype: int64
```

how many ratings at the last 5% of all ratings??

```
In [0]: print('\n No of ratings at last 5 percentile : {}'.format(sum(no_of Rated_movies_
        per_user>= 749)) )


        No of ratings at last 5 percentile : 20305
```

3.3.4 Analysis of ratings of a movie given by a user

```
In [0]: no_of_ratings_per_movie = train_df.groupby(by='movie')['rating'].count().sort_values(ascending=False)

fig = plt.figure(figsize=plt.figaspect(.5))
ax = plt.gca()
plt.plot(no_of_ratings_per_movie.values)
plt.title('# RATINGS per Movie')
plt.xlabel('Movie')
plt.ylabel('No of Users who rated a movie')
ax.set_xticklabels([])

plt.show()
```




- It is very skewed.. just like number of ratings given per user.


- There are some movies (which are very popular) which are rated by huge number of users.
- But most of the movies (like 90%) got some hundreds of ratings.

3.3.5 Number of ratings on each day of the week

```
In [0]: fig, ax = plt.subplots()
sns.countplot(x='day_of_week', data=train_df, ax=ax)
plt.title('No of ratings on each day...')
plt.ylabel('Total no of ratings')
plt.xlabel('')
ax.set_yticklabels([human(item, 'M') for item in ax.get_yticks()])
plt.show()
```



```
In [0]: start = datetime.now()
fig = plt.figure(figsize=plt.figaspect(.45))
sns.boxplot(y='rating', x='day_of_week', data=train_df)
plt.show()
print(datetime.now() - start)
```



0:01:10.003761

```
In [0]: avg_week_df = train_df.groupby(by=['day_of_week'])['rating'].mean()
print(" AVerage ratings")
print("-"*30)
print(avg_week_df)
print("\n")
```

```
AVerage ratings
-----
day_of_week
Friday      3.585274
Monday      3.577250
Saturday    3.591791
Sunday      3.594144
Thursday    3.582463
Tuesday     3.574438
Wednesday   3.583751
Name: rating, dtype: float64
```

3.3.6 Creating sparse matrix from data frame

3.3.6.1 Creating sparse matrix from train data frame


```
In [0]: start = datetime.now()
if os.path.isfile('train_sparse_matrix.npz'):
    print("It is present in your pwd, getting it from disk....")
    # just get it from the disk instead of computing it
    train_sparse_matrix = sparse.load_npz('train_sparse_matrix.npz')
    print("DONE..")
else:
    print("We are creating sparse_matrix from the dataframe..")
    # create sparse_matrix and store it for after usage.
    # csr_matrix(data_values, (row_index, col_index), shape_of_matrix)
    # It should be in such a way that, MATRIX[row, col] = data
    train_sparse_matrix = sparse.csr_matrix((train_df.rating.values, (train_df.user
    .values,
                                     train_df.movie.values)),)

    print('Done. It\'s shape is : (user, movie) : ', train_sparse_matrix.shape)
    print('Saving it into disk for furthur usage..')
    # save it into disk
    sparse.save_npz("train_sparse_matrix.npz", train_sparse_matrix)
    print('Done..\n')

print(datetime.now() - start)
```

```
It is present in your pwd, getting it from disk....
DONE..
0:00:04.463750
```

The Sparsity of Train Sparse Matrix

```
In [0]: us,mv = train_sparse_matrix.shape
elem = train_sparse_matrix.count_nonzero()

print("Sparsity Of Train matrix : {} % ".format( (1-(elem/(us*mv))) * 100) )

Sparsity Of Train matrix : 99.8292709259195 %
```

3.3.6.2 Creating sparse matrix from test data frame

```
In [0]: start = datetime.now()
if os.path.isfile('test_sparse_matrix.npz'):
    print("It is present in your pwd, getting it from disk....")
    # just get it from the disk instead of computing it
    test_sparse_matrix = sparse.load_npz('test_sparse_matrix.npz')
    print("DONE..")
else:
    print("We are creating sparse_matrix from the dataframe..")
    # create sparse_matrix and store it for after usage.
    # csr_matrix(data_values, (row_index, col_index), shape_of_matrix)
    # It should be in such a way that, MATRIX[row, col] = data
    test_sparse_matrix = sparse.csr_matrix((test_df.rating.values, (test_df.user.values,
                                                                    test_df.movie.values)))

    print('Done. It\'s shape is : (user, movie) : ', test_sparse_matrix.shape)
    print('Saving it into disk for further usage..')
    # save it into disk
    sparse.save_npz("test_sparse_matrix.npz", test_sparse_matrix)
    print('Done..\n')

print(datetime.now() - start)
```

```
It is present in your pwd, getting it from disk....
DONE..
0:00:01.172240
```

The Sparsity of Test data Matrix

```
In [0]: us, mv = test_sparse_matrix.shape
elem = test_sparse_matrix.count_nonzero()

print("Sparsity Of Test matrix : {} % ".format( (1-(elem/(us*mv))) * 100) )

Sparsity Of Test matrix : 99.95731772988694 %
```

3.3.7 Finding Global average of all movie ratings, Average rating per user, and Average rating per movie

```
In [0]: # get the user averages in dictionary (key: user_id/movie_id, value: avg rating)

def get_average_ratings(sparse_matrix, of_users):

    # average ratings of user/axes
    ax = 1 if of_users else 0 # 1 - User axes, 0 - Movie axes

    # ".A1" is for converting Column_Matrix to 1-D numpy array
    sum_of_ratings = sparse_matrix.sum(axis=ax).A1
    # Boolean matrix of ratings ( whether a user rated that movie or not)
    is_rated = sparse_matrix!=0
    # no of ratings that each user OR movie..
    no_of_ratings = is_rated.sum(axis=ax).A1

    # max_user and max_movie ids in sparse matrix
    u,m = sparse_matrix.shape
    # create a dictionary of users and their average ratings..
    average_ratings = { i : sum_of_ratings[i]/no_of_ratings[i]
                        for i in range(u if of_users else m)
                        if no_of_ratings[i] !=0}

    # return that dictionary of average ratings
    return average_ratings
```

3.3.7.1 finding global average of all movie ratings

```
In [0]: train_averages = dict()
# get the global average of ratings in our train set.
train_global_average = train_sparse_matrix.sum()/train_sparse_matrix.count_nonzero(
)
train_averages['global'] = train_global_average
train_averages
```

```
Out[0]: {'global': 3.582890686321557}
```

3.3.7.2 finding average rating per user

```
In [0]: train_averages['user'] = get_average_ratings(train_sparse_matrix, of_users=True)
print('\nAverage rating of user 10 : ',train_averages['user'][10])
```

```
Average rating of user 10 : 3.3781094527363185
```

3.3.7.3 finding average rating per movie

```
In [0]: train_averages['movie'] = get_average_ratings(train_sparse_matrix, of_users=False)
print('\nAverage rating of movie 15 : ',train_averages['movie'][15])
```

```
Average rating of movie 15 : 3.3038461538461537
```


3.3.7.4 PDF's & CDF's of Avg.Ratings of Users & Movies (In Train Data)

```
In [0]: start = datetime.now()
# draw pdfs for average rating per user and average
fig, (ax1, ax2) = plt.subplots(nrows=1, ncols=2, figsize=plt.figaspect(.5))
fig.suptitle('Avg Ratings per User and per Movie', fontsize=15)

ax1.set_title('Users-Avg-Ratings')
# get the list of average user ratings from the averages dictionary..
user_averages = [rat for rat in train_averages['user'].values()]
sns.distplot(user_averages, ax=ax1, hist=False,
              kde_kws=dict(cumulative=True), label='Cdf')
sns.distplot(user_averages, ax=ax1, hist=False, label='Pdf')

ax2.set_title('Movies-Avg-Rating')
# get the list of movie_average_ratings from the dictionary..
movie_averages = [rat for rat in train_averages['movie'].values()]
sns.distplot(movie_averages, ax=ax2, hist=False,
              kde_kws=dict(cumulative=True), label='Cdf')
sns.distplot(movie_averages, ax=ax2, hist=False, label='Pdf')

plt.show()
print(datetime.now() - start)
```



0:00:35.003443

3.3.8 Cold Start problem

3.3.8.1 Cold Start problem with Users

```
In [0]: total_users = len(np.unique(df.user))
users_train = len(train_averages['user'])
new_users = total_users - users_train

print('\nTotal number of Users :', total_users)
print('\nNumber of Users in Train data :', users_train)
print("\nNo of Users that didn't appear in train data: {} ({} %) \n ".format(new_users,
                                                                              np.round((new_users/total_users)*100, 2)))

Total number of Users : 480189

Number of Users in Train data : 405041

No of Users that didn't appear in train data: 75148 (15.65 %)
```

We might have to handle **new users (75148)** who didn't appear in train data.

3.3.8.2 Cold Start problem with Movies

```
In [0]: total_movies = len(np.unique(df.movie))
movies_train = len(train_averages['movie'])
new_movies = total_movies - movies_train

print('\nTotal number of Movies :', total_movies)
print('\nNumber of Users in Train data :', movies_train)
print("\nNo of Movies that didn't appear in train data: {} ({} %) \n ".format(new_movies,
                                                                              np.round((new_movies/total_movies)*100, 2)))

Total number of Movies : 17770

Number of Users in Train data : 17424

No of Movies that didn't appear in train data: 346 (1.95 %)
```

We might have to handle **346 movies** (small comparatively) in test data

3.4 Computing Similarity matrices

3.4.1 Computing User-User Similarity matrix

1. Calculating User User Similarity_Matrix is **not very easy**(*unless you have huge Computing Power and lots of time*) because of number of. usersbeing lare.
 - You can try if you want to. Your system could crash or the program stops with **Memory Error**

3.4.1.1 Trying with all dimensions (17k dimensions per user)

```

In [0]: from sklearn.metrics.pairwise import cosine_similarity

def compute_user_similarity(sparse_matrix, compute_for_few=False, top = 100, verbose=False, verb_for_n_rows = 20,
                           draw_time_taken=True):
    no_of_users, _ = sparse_matrix.shape
    # get the indices of non zero rows(users) from our sparse matrix
    row_ind, col_ind = sparse_matrix.nonzero()
    row_ind = sorted(set(row_ind)) # we don't have to
    time_taken = list() # time taken for finding similar users for an user..

    # we create rows, cols, and data lists.., which can be used to create sparse matrices
    rows, cols, data = list(), list(), list()
    if verbose: print("Computing top",top,"similarities for each user..")

    start = datetime.now()
    temp = 0

    for row in row_ind[:top] if compute_for_few else row_ind:
        temp = temp+1
        prev = datetime.now()

        # get the similarity row for this user with all other users
        sim = cosine_similarity(sparse_matrix.getrow(row), sparse_matrix).ravel()
        # We will get only the top 'top' most similar users and ignore rest of them..

        top_sim_ind = sim.argsort()[-top:]
        top_sim_val = sim[top_sim_ind]

        # add them to our rows, cols and data
        rows.extend([row]*top)
        cols.extend(top_sim_ind)
        data.extend(top_sim_val)
        time_taken.append(datetime.now().timestamp() - prev.timestamp())
        if verbose:
            if temp%verb_for_n_rows == 0:
                print("computing done for {} users [ time elapsed : {} ]".format(temp, datetime.now()-start))

    # lets create sparse matrix out of these and return it
    if verbose: print('Creating Sparse matrix from the computed similarities')
    #return rows, cols, data

    if draw_time_taken:
        plt.plot(time_taken, label = 'time taken for each user')
        plt.plot(np.cumsum(time_taken), label='Total time')
        plt.legend(loc='best')
        plt.xlabel('User')
        plt.ylabel('Time (seconds)')
        plt.show()


    return sparse.csr_matrix((data, (rows, cols)), shape=(no_of_users, no_of_users), time_taken)

```

```
In [0]: start = datetime.now()
u_u_sim_sparse, _ = compute_user_similarity(train_sparse_matrix, compute_for_few=True,
ue, top = 100,
                                             verbose=True)

print("-"*100)
print("Time taken :",datetime.now()-start)
```

```
Computing top 100 similarities for each user..
computing done for 20 users [ time elapsed : 0:03:20.300488 ]
computing done for 40 users [ time elapsed : 0:06:38.518391 ]
computing done for 60 users [ time elapsed : 0:09:53.143126 ]
computing done for 80 users [ time elapsed : 0:13:10.080447 ]
computing done for 100 users [ time elapsed : 0:16:24.711032 ]
Creating Sparse matrix from the computed similarities
```



```
-----
Time taken : 0:16:33.618931
```

3.4.1.2 Trying with reduced dimensions (Using TruncatedSVD for dimensionality reduction of user vector)

- We have **405,041 users** in our training set and computing similarities between them..(**17K dimensional vector**..) is time consuming..
- From above plot, It took roughly **8.88 sec** for computing similar users for **one user**
- We have **405,041 users** with us in training set.
- $405041 \times 8.88 = 3596764.08 \text{ sec} = 59946.068 \text{ min} = 999.101133333 \text{ hours} = 41.629213889 \text{ days} \dots$
 - Even if we run on 4 cores parallelly (a typical system now a days), It will still take almost **10 and 1/2 days**.

IDEA: Instead, we will try to reduce the dimensions using SVD, so that **it might** speed up the process...

```
In [0]: from datetime import datetime
        from sklearn.decomposition import TruncatedSVD

        start = datetime.now()

        # initialize the algorithm with some parameters..
        # All of them are default except n_components. n_iter is for Randomized SVD solver.
        netflix_svd = TruncatedSVD(n_components=500, algorithm='randomized', random_state=15)
        trunc_svd = netflix_svd.fit_transform(train_sparse_matrix)

        print(datetime.now()-start)

0:29:07.069783
```

Here,

- $\Sigma \leftarrow (\text{netflix_svd.singular_values_})$
- $V^T \leftarrow (\text{netflix_svd.components_})$
- U is not returned. instead **Projection_of_X** onto the new vectorspace is returned.
- It uses **randomized svd** internally, which returns **All 3 of them separately**. Use that instead..

```
In [0]: expl_var = np.cumsum(netflix_svd.explained_variance_ratio_)
```


```
In [0]: fig, (ax1, ax2) = plt.subplots(nrows=1, ncols=2, figsize=plt.figaspect(.5))

ax1.set_ylabel("Variance Explained", fontsize=15)
ax1.set_xlabel("# Latent Facors", fontsize=15)
ax1.plot(expl_var)
# annotate some (latentfactors, expl_var) to make it clear
ind = [1, 2, 4, 8, 20, 60, 100, 200, 300, 400, 500]
ax1.scatter(x = [i-1 for i in ind], y = expl_var[[i-1 for i in ind]], c='#ff3300')
for i in ind:
    ax1.annotate(s = "({}, {})".format(i, np.round(expl_var[i-1], 2)), xy=(i-1, expl_var[i-1]),
                xytext = (i+20, expl_var[i-1] - 0.01), fontweight='bold')

change_in_expl_var = [expl_var[i+1] - expl_var[i] for i in range(len(expl_var)-1)]
ax2.plot(change_in_expl_var)

ax2.set_ylabel("Gain in Var_Expl with One Additional LF", fontsize=10)
ax2.yaxis.set_label_position("right")
ax2.set_xlabel("# Latent Facors", fontsize=20)

plt.show()
```



```
In [0]: for i in ind:
        print("({}, {})".format(i, np.round(expl_var[i-1], 2)))

(1, 0.23)
(2, 0.26)
(4, 0.3)
(8, 0.34)
(20, 0.38)
(60, 0.44)
(100, 0.47)
(200, 0.53)
(300, 0.57)
(400, 0.61)
(500, 0.64)
```

I think 500 dimensions is good enough

- By just taking **(20 to 30)** latent factors, explained variance that we could get is **20 %**.
- To take it to **60%**, we have to take **almost 400 latent factors**. It is not fare.
- It basically is the **gain of variance explained**, if we **add one additional latent factor to it**.
- By adding one by one latent factor too it, the **_gain in explained variance** with that addition is decreasing. (Obviously, because they are sorted that way).
- **LHS Graph:**
 - **x** --- (No of latent factos),
 - **y** --- (The variance explained by taking x latent factors)
- **More decrease in the line (RHS graph) :**
 - We are getting more explained variance than before.
- **Less decrease in that line (RHS graph) :**
 - We are not getting benifitted from adding latent factor furthur. This is what is shown in the plots.
- **RHS Graph:**
 - **x** --- (No of latent factors),
 - **y** --- (Gain n Expl_Var by taking one additional latent factor)

```
In [0]: # Let's project our Original U_M matrix into into 500 Dimensional space...
start = datetime.now()
trunc_matrix = train_sparse_matrix.dot(netflix_svd.components_.T)
print(datetime.now() - start)

0:00:45.670265
```

```
In [0]: type(trunc_matrix), trunc_matrix.shape
```

```
Out[0]: (numpy.ndarray, (2649430, 500))
```

- Let's convert this to actual sparse matrix and store it for future purposes

```
In [0]: if not os.path.isfile('trunc_sparse_matrix.npz'):
# create that sparse sparse matrix
trunc_sparse_matrix = sparse.csr_matrix(trunc_matrix)
# Save this truncated sparse matrix for later usage..
sparse.save_npz('trunc_sparse_matrix', trunc_sparse_matrix)
else:
trunc_sparse_matrix = sparse.load_npz('trunc_sparse_matrix.npz')
```


```
In [0]: trunc_sparse_matrix.shape
```

```
Out[0]: (2649430, 500)
```

```
In [0]: start = datetime.now()
trunc_u_u_sim_matrix, _ = compute_user_similarity(trunc_sparse_matrix, compute_for_
few=True, top=50, verbose=True,
                                                    verb_for_n_rows=10)

print("-"*50)
print("time:", datetime.now()-start)
```

```
Computing top 50 similarities for each user..
computing done for 10 users [ time elapsed : 0:02:09.746324 ]
computing done for 20 users [ time elapsed : 0:04:16.017768 ]
computing done for 30 users [ time elapsed : 0:06:20.861163 ]
computing done for 40 users [ time elapsed : 0:08:24.933316 ]
computing done for 50 users [ time elapsed : 0:10:28.861485 ]
Creating Sparse matrix from the computed similarities
```



```
-----
time: 0:10:52.658092
```

: This is taking more time for each user than Original one.

- from above plot, It took almost **12.18** for computing similar users for **one user**
- We have **405041 users** with us in training set.
- $405041 \times 12.18 = 4933399.38 \text{ sec} = 82223.323 \text{ min} = 1370.388716667 \text{ hours} = 57.09$
 - Even we run on 4 cores parallelly (a typical system now a days), It will still take almost **(14 - 15)** days.

- Why did this happen...??

- Just think about it. It's not that difficult.

-----(*sparse & dense.....get it ??*)-----

Is there any other way to compute user user similarity..??

-An alternative is to compute similar users for a particular user, whenever required (**ie., Run time**)

- We maintain a binary Vector for users, which tells us whether we already computed or not..
- *****If not*** :**
 - Compute top (let's just say, 1000) most similar users for this given user, and add this to our datastructure, so that we can just access it(similar users) without recomputing it again.
-
- *****If It is already Computed***:**
 - Just get it directly from our datastructure, which has that information.
 - In production time, We might have to recompute similarities, if it is computed a long time ago. Because user preferences changes over time. If we could maintain some kind of Timer, which when expires, we have to update it (recompute it).
-
- *****Which datastructure to use:*****
 - It is purely implementation dependant.
 - One simple method is to maintain a ****Dictionary Of Dictionaries****.
 -
 - ****key : _userid_**
 - **__value__ : _Again a dictionary_**
 - **__key__ : _Similar User_**
 - **__value__ : _Similarity Value_**

3.4.2 Computing Movie-Movie Similarity matrix

```
In [0]: start = datetime.now()
if not os.path.isfile('m_m_sim_sparse.npz'):
    print("It seems you don't have that file. Computing movie_movie similarity...")
    start = datetime.now()
    m_m_sim_sparse = cosine_similarity(X=train_sparse_matrix.T, dense_output=False)
    print("Done..")
    # store this sparse matrix in disk before using it. For future purposes.
    print("Saving it to disk without the need of re-computing it again.. ")
    sparse.save_npz("m_m_sim_sparse.npz", m_m_sim_sparse)
    print("Done..")
else:
    print("It is there, We will get it.")
    m_m_sim_sparse = sparse.load_npz("m_m_sim_sparse.npz")
    print("Done ...")

print("It's a ", m_m_sim_sparse.shape, " dimensional matrix")

print(datetime.now() - start)
```

```
It seems you don't have that file. Computing movie_movie similarity...
Done..
Saving it to disk without the need of re-computing it again..
Done..
It's a (17771, 17771) dimensional matrix
0:10:02.736054
```

```
In [0]: m_m_sim_sparse.shape
```

```
Out[0]: (17771, 17771)
```

- Even though we have similarity measure of each movie, with all other movies, We generally don't care much about least similar movies.
- Most of the times, only top_xxx similar items matters. It may be 10 or 100.
- We take only those top similar movie ratings and store them in a saperate dictionary.

```
In [0]: movie_ids = np.unique(m_m_sim_sparse.nonzero()[1])
```

```
In [0]: start = datetime.now()
similar_movies = dict()
for movie in movie_ids:
    # get the top similar movies and store them in the dictionary
    sim_movies = m_m_sim_sparse[movie].toarray().ravel().argsort()[::-1][1:]
    similar_movies[movie] = sim_movies[:100]
print(datetime.now() - start)

# just testing similar movies for movie_15
similar_movies[15]
```

0:00:33.411700

```
Out[0]: array([ 8279,  8013, 16528,  5927, 13105, 12049,  4424, 10193, 17590,
                4549,  3755,   590, 14059, 15144, 15054,  9584,  9071,  6349,
                16402,  3973,  1720,  5370, 16309,  9376,  6116,  4706,  2818,
                 778, 15331,  1416, 12979, 17139, 17710,  5452,  2534,   164,
                15188,  8323,  2450, 16331,  9566, 15301, 13213, 14308, 15984,
                10597,  6426,  5500,  7068,  7328,  5720,  9802,   376, 13013,
                 8003, 10199,  3338, 15390,  9688, 16455, 11730,  4513,   598,
                12762,  2187,   509,  5865,  9166, 17115, 16334,  1942,  7282,
                17584,  4376,  8988,  8873,  5921,  2716, 14679, 11947, 11981,
                 4649,   565, 12954, 10788, 10220, 10963,  9427,  1690,  5107,
                 7859,  5969,  1510,  2429,   847,  7845,  6410, 13931,  9840,
                3706])
```

3.4.3 Finding most similar movies using similarity matrix

Does Similarity really works as the way we expected...?

Let's pick some random movie and check for its similar movies....

```
In [0]: # First Let's load the movie details into soe dataframe..
# movie details are in 'netflix/movie_titles.csv'

movie_titles = pd.read_csv("data_folder/movie_titles.csv", sep=',', header = None,
                           names=['movie_id', 'year_of_release', 'title'], verbose=
True,
                           index_col = 'movie_id', encoding = "ISO-8859-1")

movie_titles.head()
```

Tokenization took: 4.50 ms
Type conversion took: 165.72 ms
Parser memory cleanup took: 0.01 ms

Out[0]:

	year_of_release	title
movie_id		
1	2003.0	Dinosaur Planet
2	2004.0	Isle of Man TT 2004 Review
3	1997.0	Character
4	1994.0	Paula Abdul's Get Up & Dance
5	2004.0	The Rise and Fall of ECW

Similar Movies for 'Vampire Journals'

```
In [0]: mv_id = 67

print("\nMovie ----->", movie_titles.loc[mv_id].values[1])

print("\nIt has {} Ratings from users.".format(train_sparse_matrix[:, mv_id].getnnz(
)))

print("\nWe have {} movies which are similar to this and we will get only top most.
.".format(m_m_sim_sparse[:, mv_id].getnnz()))
```

Movie -----> Vampire Journals

It has 270 Ratings from users.

We have 17284 movies which are similar to this and we will get only top most..

```
In [0]: similarities = m_m_sim_sparse[mv_id].toarray().ravel()

similar_indices = similarities.argsort()[::-1][1:]


similarities[similar_indices]

sim_indices = similarities.argsort()[::-1][1:] # It will sort and reverse the array
and ignore its similarity (ie., 1)

# and return its indices (movie_ids)
```



```
In [0]: plt.plot(similarities[sim_indices], label='All the ratings')
plt.plot(similarities[sim_indices[:100]], label='top 100 similar movies')
plt.title("Similar Movies of {}".format(mv_id), fontsize=20)
plt.xlabel("Movies (Not Movie_Ids)", fontsize=15)
plt.ylabel("Cosine Similarity", fontsize=15)
plt.legend()
plt.show()
```



Top 10 similar movies

```
In [0]: movie_titles.loc[sim_indices[:10]]
```

```
Out[0]:
```

	year_of_release	title
movie_id		
323	1999.0	Modern Vampires
4044	1998.0	Subspecies 4: Bloodstorm
1688	1993.0	To Sleep With a Vampire
13962	2001.0	Dracula: The Dark Prince
12053	1993.0	Dracula Rising
16279	2002.0	Vampires: Los Muertos
4667	1996.0	Vampirella
1900	1997.0	Club Vampire
13873	2001.0	The Breed
15867	2003.0	Dracula II: Ascension

Similarly, we can *find similar users* and compare how similar they are.

4. Machine Learning Models

```

In [0]: def get_sample_sparse_matrix(sparse_matrix, no_users, no_movies, path, verbose = True):
    """
        It will get it from the 'path' if it is present or It will create
        and store the sampled sparse matrix in the path specified.
    """

    # get (row, col) and (rating) tuple from sparse_matrix...
    row_ind, col_ind, ratings = sparse.find(sparse_matrix)
    users = np.unique(row_ind)
    movies = np.unique(col_ind)

    print("Original Matrix : (users, movies) -- ({}) {}".format(len(users), len(movies)))
    print("Original Matrix : Ratings -- {}\n".format(len(ratings)))

    # It just to make sure to get same sample everytime we run this program..
    # and pick without replacement....
    np.random.seed(15)
    sample_users = np.random.choice(users, no_users, replace=False)
    sample_movies = np.random.choice(movies, no_movies, replace=False)
    # get the boolean mask or these sampled items in original row/col inds..
    mask = np.logical_and( np.isin(row_ind, sample_users),
                           np.isin(col_ind, sample_movies) )

    sample_sparse_matrix = sparse.csr_matrix((ratings[mask], (row_ind[mask], col_ind[mask])),
                                              shape=(max(sample_users)+1, max(sample_movies)+1))

    if verbose:
        print("Sampled Matrix : (users, movies) -- ({}) {}".format(len(sample_users), len(sample_movies)))
        print("Sampled Matrix : Ratings --", format(ratings[mask].shape[0]))

    print('Saving it into disk for further usage..')
    # save it into disk
    sparse.save_npz(path, sample_sparse_matrix)
    if verbose:
        print('Done..\n')

    return sample_sparse_matrix

```

4.1 Sampling Data

4.1.1 Build sample train data from the train data

```
In [0]: start = datetime.now()
path = "sample/small/sample_train_sparse_matrix25.npz"
if os.path.isfile(path):
    print("It is present in your pwd, getting it from disk....")
    # just get it from the disk instead of computing it
    sample_train_sparse_matrix = sparse.load_npz(path)
    print("DONE..")
else:
    # get 10k users and 1k movies from available data
    sample_train_sparse_matrix = get_sample_sparse_matrix(train_sparse_matrix, no_u
sers=18000, no_movies=3000,
                                                         path = path)

print(datetime.now() - start)
```

Original Matrix : (users, movies) -- (405041 17424)
Original Matrix : Ratings -- 80384405

Sampled Matrix : (users, movies) -- (18000 3000)
Sampled Matrix : Ratings -- 617650
Saving it into disk for furthur usage..
Done..

0:01:33.017998

4.1.2 Build sample test data from the test data

```
In [0]: start = datetime.now()

path = "sample/small/sample_test_sparse_matrix.npz"
if os.path.isfile(path):
    print("It is present in your pwd, getting it from disk....")
    # just get it from the disk instead of computing it
    sample_test_sparse_matrix = sparse.load_npz(path)
    print("DONE..")
else:
    # get 5k users and 500 movies from available data
    sample_test_sparse_matrix = get_sample_sparse_matrix(test_sparse_matrix, no_use
rs=9000, no_movies=1500,
                                                         path = "sample/small/sample_test_s
parse_matrix.npz")
print(datetime.now() - start)
```

It is present in your pwd, getting it from disk....
DONE..
0:00:02.501058

4.2 Finding Global Average of all movie ratings, Average rating per User, and Average rating per Movie (from sampled train)

```
In [0]: sample_train_averages = dict()
```

4.2.1 Finding Global Average of all movie ratings

```
In [0]: # get the global average of ratings in our train set.
global_average = sample_train_sparse_matrix.sum()/sample_train_sparse_matrix.count_nonzero()
sample_train_averages['global'] = global_average
sample_train_averages
```

```
Out[0]: {'global': 3.5869877762486846}
```

4.2.2 Finding Average rating per User

```
In [0]: sample_train_averages['user'] = get_average_ratings(sample_train_sparse_matrix, of_
users=True)
print('\nAverage rating of user 1515220 : ',sample_train_averages['user'][1515220])
```

```
Average rating of user 1515220 : 3.923076923076923
```

4.2.3 Finding Average rating per Movie

```
In [0]: sample_train_averages['movie'] = get_average_ratings(sample_train_sparse_matrix, o
f_users=False)
print('\n Average rating of movie 15153 : ',sample_train_averages['movie'][15153])
```

```
Average rating of movie 15153 : 2.6555555555555554
```

4.3 Featurizing data

```
In [0]: print('\n No of ratings in Our Sampled train matrix is : {}'.format(sample_train_
sparse_matrix.count_nonzero()))
print('\n No of ratings in Our Sampled test  matrix is : {}'.format(sample_test_s
parse_matrix.count_nonzero()))
```

```
No of ratings in Our Sampled train matrix is : 617650
```

```
No of ratings in Our Sampled test  matrix is : 48930
```

4.3.1 Featurizing data for regression problem

4.3.1.1 Featurizing train data

```
In [0]: # get users, movies and ratings from our samples train sparse matrix
sample_train_users, sample_train_movies, sample_train_ratings = sparse.find(sample_
train_sparse_matrix)
```

```

In [0]: #####
# It took me almost 10 hours to prepare this train dataset.#
#####
start = datetime.now()
if os.path.isfile('sample/small/reg_train2.csv'):
    print("File already exists you don't have to prepare again..." )
else:
    print('preparing {} tuples for the dataset..\\n'.format(len(sample_train_ratings)
    ))
    with open('sample/small/reg_train.csv', mode='w') as reg_data_file:
        count = 0
        for (user, movie, rating) in zip(sample_train_users, sample_train_movies,
sample_train_ratings):
            st = datetime.now()
            #         print(user, movie)
            #----- Ratings of "movie" by similar users of "user" --
            -----
            # compute the similar Users of the "user"
            user_sim = cosine_similarity(sample_train_sparse_matrix[user], sample_t
rain_sparse_matrix).ravel()
            top_sim_users = user_sim.argsort()[::-1][1:] # we are ignoring 'The Use
r' from its similar users.
            # get the ratings of most similar users for this movie
            top_ratings = sample_train_sparse_matrix[top_sim_users, movie].toarray(
).ravel()
            # we will make it's length "5" by adding movie averages to .
            top_sim_users_ratings = list(top_ratings[top_ratings != 0][:5])
            top_sim_users_ratings.extend([sample_train_averages['movie'][movie]]*(5
- len(top_sim_users_ratings)))
            #         print(top_sim_users_ratings, end=" ")

            #----- Ratings by "user" to similar movies of "movie"
            -----
            # compute the similar movies of the "movie"
            movie_sim = cosine_similarity(sample_train_sparse_matrix[:,movie].T, sa
mple_train_sparse_matrix.T).ravel()
            top_sim_movies = movie_sim.argsort()[::-1][1:] # we are ignoring 'The U
ser' from its similar users.
            # get the ratings of most similar movie rated by this user..
            top_ratings = sample_train_sparse_matrix[user, top_sim_movies].toarray(
).ravel()
            # we will make it's length "5" by adding user averages to.
            top_sim_movies_ratings = list(top_ratings[top_ratings != 0][:5])
            top_sim_movies_ratings.extend([sample_train_averages['user'][user]]*(5-
len(top_sim_movies_ratings)))
            #         print(top_sim_movies_ratings, end=" : -- ")

            #-----prepare the row to be stores in a file-----
--#
            row = list()
            row.append(user)
            row.append(movie)
            # Now add the other features to this data...
            row.append(sample_train_averages['global']) # first feature
            # next 5 features are similar_users "movie" ratings
            row.extend(top_sim_users_ratings)
            # next 5 features are "user" ratings for similar_movies
            row.extend(top_sim_movies_ratings)
            # Avg_user rating
            row.append(sample_train_averages['user'][user])
            # Avg_movie rating
            row.append(sample_train_averages['movie'][movie])

```

File already exists you don't have to prepare again...
0:00:00.328261

Reading from the file to make a Train_dataframe

```
In [0]: reg_train = pd.read_csv('sample/small/reg_train2.csv', names = ['user', 'movie', 'GAvg', 'sur1', 'sur2', 'sur3', 'sur4', 'sur5', 'smr1', 'smr2', 'smr3', 'smr4', 'smr5', 'UAvg', 'MAvg', 'rating'], header=None)
reg_train.head()
```

```
Out[0]:
```

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	smr3	smr4	smr5	UAvg	MAvg
0	174683	10	3.586988	5.0	4.0	4.0	3.0	4.0	3.0	5.0	3.0	3.0	2.0	3.882353	3.636364
1	233949	10	3.586988	4.0	4.0	5.0	1.0	5.0	2.0	3.0	3.0	3.0	3.0	2.692308	3.636364
2	767518	10	3.586988	5.0	4.0	4.0	4.0	3.0	5.0	5.0	4.0	4.0	4.0	3.884615	3.636364
3	894393	10	3.586988	3.0	5.0	4.0	5.0	5.0	4.0	4.0	4.0	4.0	4.0	4.000000	3.636364
4	951907	10	3.586988	5.0	4.0	3.0	4.0	5.0	3.0	4.0	4.0	4.0	3.0	3.881188	3.636364

- **GAvg** : Average rating of all the ratings
- **Similar users rating of this movie:**
 - sur1, sur2, sur3, sur4, sur5 (top 5 similar users who rated that movie..)
- **Similar movies rated by this user:**
 - smr1, smr2, smr3, smr4, smr5 (top 5 similar movies rated by this movie..)
- **UAvg** : User's Average rating
- **MAvg** : Average rating of this movie
- **rating** : Rating of this movie by this user.

4.3.1.2 Featurizing test data

```
In [0]: # get users, movies and ratings from the Sampled Test
sample_test_users, sample_test_movies, sample_test_ratings = sparse.find(sample_test_sparse_matrix)
```

```
In [0]: sample_train_averages['global']
```

```
Out[0]: 3.5869877762486846
```

```

In [0]: start = datetime.now()

if os.path.isfile('sample/small/reg_test.csv'):
    print("It is already created...")
else:

    print('preparing {} tuples for the dataset...\n'.format(len(sample_test_ratings)
))
    with open('sample/small/reg_test.csv', mode='w') as reg_data_file:
        count = 0
        for (user, movie, rating) in zip(sample_test_users, sample_test_movies, sa
mple_test_ratings):
            st = datetime.now()

            #----- Ratings of "movie" by similar users of "user" -----
            -----
            #print(user, movie)
            try:
                # compute the similar Users of the "user"
                user_sim = cosine_similarity(sample_train_sparse_matrix[user], samp
le_train_sparse_matrix).ravel()
                top_sim_users = user_sim.argsort()[::-1][1:] # we are ignoring 'The
User' from its similar users.
                # get the ratings of most similar users for this movie
                top_ratings = sample_train_sparse_matrix[top_sim_users, movie].toar
ray().ravel()

                # we will make it's length "5" by adding movie averages to .
                top_sim_users_ratings = list(top_ratings[top_ratings != 0][:5])
                top_sim_users_ratings.extend([sample_train_averages['movie'][movie]
]*(5 - len(top_sim_users_ratings)))
                # print(top_sim_users_ratings, end="--")

            except (IndexError, KeyError):
                # It is a new User or new Movie or there are no ratings for given u
ser for top similar movies...
                ##### Cold Start Problem #####
                top_sim_users_ratings.extend([sample_train_averages['global']]*(5 -
len(top_sim_users_ratings)))
                #print(top_sim_users_ratings)
            except:
                print(user, movie)
                # we just want KeyErrors to be resolved. Not every Exception...
                raise

            #----- Ratings by "user" to similar movies of "movie" -----
            -----
            try:
                # compute the similar movies of the "movie"
                movie_sim = cosine_similarity(sample_train_sparse_matrix[:,movie].T
, sample_train_sparse_matrix.T).ravel()
                top_sim_movies = movie_sim.argsort()[::-1][1:] # we are ignoring 'T
he User' from its similar users.
                # get the ratings of most similar movie rated by this user..
                top_ratings = sample_train_sparse_matrix[user, top_sim_movies].toar
ray().ravel()

                # we will make it's length "5" by adding user averages to.
                top_sim_movies_ratings = list(top_ratings[top_ratings != 0][:5])
                top_sim_movies_ratings.extend([sample_train_averages['user'][user]]
*(5-len(top_sim_movies_ratings)))
                #print(top_sim_movies_ratings)
            except (IndexError, KeyError):
                #print(top sim movies ratings. end=" : -- ")

```


It is already created...

Reading from the file to make a test dataframe

```
In [0]: reg_test_df = pd.read_csv('sample/small/reg_test.csv', names = ['user', 'movie', 'GAvg', 'sur1', 'sur2', 'sur3', 'sur4', 'sur5', 'smr1', 'smr2', 'smr3', 'smr4', 'smr5', 'UAvg', 'MAvg', 'rating'], header=None)
reg_test_df.head(4)
```

Out[0]:

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	smr3
0	808635	71	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988
1	898730	71	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988
2	941866	71	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988
3	1280761	71	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988

- **GAvg** : Average rating of all the ratings
- **Similar users rating of this movie:**
 - sur1, sur2, sur3, sur4, sur5 (top 5 simiular users who rated that movie..)
- **Similar movies rated by this user:**
 - smr1, smr2, smr3, smr4, smr5 (top 5 simiular movies rated by this movie..)
- **UAvg** : User AVerage rating
- **MAvg** : Average rating of this movie
- **rating** : Rating of this movie by this user.

4.3.2 Transforming data for Surprise models

```
In [0]: from surprise import Reader, Dataset
```

4.3.2.1 Transforming train data

- We can't give raw data (movie, user, rating) to train the model in Surprise library.
- They have a separate format for TRAIN and TEST data, which will be useful for training the models like SVD, KNNBaseLineOnly....etc., in Surprise.
- We can form the trainset from a file, or from a Pandas DataFrame. http://surprise.readthedocs.io/en/stable/getting_started.html#load-dom-dataframe-py (http://surprise.readthedocs.io/en/stable/getting_started.html#load-dom-dataframe-py)

```
In [0]: # It is to specify how to read the dataframe.
# for our dataframe, we don't have to specify anything extra..
reader = Reader(rating_scale=(1,5))

# create the traindata from the dataframe...
train_data = Dataset.load_from_df(reg_train[['user', 'movie', 'rating']], reader)

# build the trainset from traindata..., It is of dataset format from surprise library..
trainset = train_data.build_full_trainset()
```

4.3.2.2 Transforming test data

- Testset is just a list of (user, movie, rating) tuples. (Order in the tuple is important)

```
In [0]: testset = list(zip(reg_test_df.user.values, reg_test_df.movie.values, reg_test_df.rating.values))
testset[:3]
```

```
Out[0]: [(808635, 71, 5), (898730, 71, 3), (941866, 71, 4)]
```

4.4 Applying Machine Learning models

- Global dictionary that stores rmse and mape for all the models....
 - It stores the metrics in a dictionary of dictionaries

```
keys : model names(string)
value: dict(key : metric, value : value )
```

```
In [0]: models_evaluation_train = dict()
models_evaluation_test = dict()

models_evaluation_train, models_evaluation_test
```

```
Out[0]: ({}, {})
```

Utility functions for running regression models

```

In [0]: # to get rmse and mape given actual and predicted ratings..
def get_error_metrics(y_true, y_pred):
    rmse = np.sqrt(np.mean([ (y_true[i] - y_pred[i])**2 for i in range(len(y_pred))
]))
    mape = np.mean(np.abs( (y_true - y_pred)/y_true )) * 100
    return rmse, mape

#####
#####
def run_xgboost(algo, x_train, y_train, x_test, y_test, verbose=True):
    """
    It will return train_results and test_results
    """

    # dictionaries for storing train and test results
    train_results = dict()
    test_results = dict()

    # fit the model
    print('Training the model..')
    start = datetime.now()
    algo.fit(x_train, y_train, eval_metric = 'rmse')
    print('Done. Time taken : {}'.format(datetime.now()-start))
    print('Done \n')

    # from the trained model, get the predictions....
    print('Evaluating the model with TRAIN data...')
    start = datetime.now()
    y_train_pred = algo.predict(x_train)
    # get the rmse and mape of train data...
    rmse_train, mape_train = get_error_metrics(y_train.values, y_train_pred)

    # store the results in train_results dictionary..
    train_results = {'rmse': rmse_train,
                    'mape' : mape_train,
                    'predictions' : y_train_pred}

    #####
    # get the test data predictions and compute rmse and mape
    print('Evaluating Test data')
    y_test_pred = algo.predict(x_test)
    rmse_test, mape_test = get_error_metrics(y_true=y_test.values, y_pred=y_test_pr
ed)

    # store them in our test results dictionary.
    test_results = {'rmse': rmse_test,
                  'mape' : mape_test,
                  'predictions':y_test_pred}

    if verbose:
        print('\nTEST DATA')
        print('-'*30)
        print('RMSE : ', rmse_test)
        print('MAPE : ', mape_test)

    # return these train and test results...
    return train_results, test_results

```

Utility functions for Surprise modes

```

In [0]: # it is just to make sure that all of our algorithms should produce same results
        # everytime they run...

my_seed = 15
random.seed(my_seed)
np.random.seed(my_seed)

#####
# get (actual_list , predicted_list) ratings given list
# of predictions (prediction is a class in Surprise).
#####
def get_ratings(predictions):
    actual = np.array([pred.r_ui for pred in predictions])
    pred = np.array([pred.est for pred in predictions])

    return actual, pred

#####
# get 'rmse' and 'mape' , given list of prediction objects
#####
def get_errors(predictions, print_them=False):

    actual, pred = get_ratings(predictions)
    rmse = np.sqrt(np.mean((pred - actual)**2))
    mape = np.mean(np.abs(pred - actual)/actual)

    return rmse, mape*100

#####
# It will return predicted ratings, rmse and mape of both train and test data #
#####
def run_surprise(algo, trainset, testset, verbose=True):
    '''
        return train_dict, test_dict

        It returns two dictionaries, one for train and the other is for test
        Each of them have 3 key-value pairs, which specify 'rmse', 'mape', and
        'predicted ratings'.
    '''
    start = datetime.now()
    # dictionaries that stores metrics for train and test..
    train = dict()
    test = dict()

    # train the algorithm with the trainset
    st = datetime.now()
    print('Training the model...')
    algo.fit(trainset)
    print('Done. time taken : {} \n'.format(datetime.now()-st))

    # ----- Evaluating train data-----#
    st = datetime.now()
    print('Evaluating the model with train data..')
    # get the train predictions (list of prediction class inside Surprise)
    train_preds = algo.test(trainset.build_testset())
    # get predicted ratings from the train predictions..
    train_actual_ratings, train_pred_ratings = get_ratings(train_preds)
    # get 'rmse' and 'mape' from the train predictions.
    train_rmse, train_mape = get_errors(train_preds)
    print('time taken : {}'.format(datetime.now()-st))

    if verbose:
        print('-'*15)
        print('Train Data')

```

4.4.1 XGBoost with initial 13 features

```
In [0]: import xgboost as xgb
```

```
In [0]: parameters = {'max_depth': [1, 2, 3],  
                      'learning_rate': [0.001, 0.01, 0.1],  
                      'n_estimators': [100, 300, 500, 700]}
```

```
In [0]: from sklearn.model_selection import GridSearchCV  
        from sklearn.model_selection import TimeSeriesSplit
```

```
In [0]: # prepare Train data
x_train = reg_train.drop(['user', 'movie', 'rating'], axis=1)
y_train = reg_train['rating']

# Prepare Test data
x_test = reg_test_df.drop(['user', 'movie', 'rating'], axis=1)
y_test = reg_test_df['rating']

start = datetime.now()
# initialize Our first XGBoost model...
first_xgb = xgb.XGBRegressor(nthread=-1)

# Perform cross validation
gscv = GridSearchCV(first_xgb,
                    param_grid = parameters,
                    scoring="neg_mean_squared_error",
                    cv = TimeSeriesSplit(n_splits=5),
                    n_jobs = -1,
                    verbose = 1)
gscv_result = gscv.fit(x_train, y_train)

# Summarize results
print("Best: %f using %s" % (gscv_result.best_score_, gscv_result.best_params_))
means = gscv_result.cv_results_['mean_test_score']
stds = gscv_result.cv_results_['std_test_score']
params = gscv_result.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))

print("\nTime Taken: ", start - datetime.now())
```


Fitting 5 folds for each of 36 candidates, totalling 180 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.  
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed: 19.3min  
[Parallel(n_jobs=-1)]: Done 180 out of 180 | elapsed: 100.8min finished
```

```
Best: -0.744863 using {'learning_rate': 0.1, 'max_depth': 3, 'n_estimators': 500
}
-8.940064 (0.136602) with: {'learning_rate': 0.001, 'max_depth': 1, 'n_estimator
s': 100}
-6.322584 (0.104569) with: {'learning_rate': 0.001, 'max_depth': 1, 'n_estimator
s': 300}
-4.560439 (0.079736) with: {'learning_rate': 0.001, 'max_depth': 1, 'n_estimator
s': 500}
-3.375169 (0.062092) with: {'learning_rate': 0.001, 'max_depth': 1, 'n_estimator
s': 700}
-8.910045 (0.119083) with: {'learning_rate': 0.001, 'max_depth': 2, 'n_estimator
s': 100}
-6.268585 (0.086633) with: {'learning_rate': 0.001, 'max_depth': 2, 'n_estimator
s': 300}
-4.494120 (0.066408) with: {'learning_rate': 0.001, 'max_depth': 2, 'n_estimator
s': 500}
-3.303398 (0.052270) with: {'learning_rate': 0.001, 'max_depth': 2, 'n_estimator
s': 700}
-8.900899 (0.118331) with: {'learning_rate': 0.001, 'max_depth': 3, 'n_estimator
s': 100}
-6.245090 (0.084058) with: {'learning_rate': 0.001, 'max_depth': 3, 'n_estimator
s': 300}
-4.463697 (0.061850) with: {'learning_rate': 0.001, 'max_depth': 3, 'n_estimator
s': 500}
-3.267050 (0.045774) with: {'learning_rate': 0.001, 'max_depth': 3, 'n_estimator
s': 700}
-2.271043 (0.045529) with: {'learning_rate': 0.01, 'max_depth': 1, 'n_estimators
': 100}
-0.897886 (0.022963) with: {'learning_rate': 0.01, 'max_depth': 1, 'n_estimators
': 300}
-0.819336 (0.023904) with: {'learning_rate': 0.01, 'max_depth': 1, 'n_estimators
': 500}
-0.790015 (0.023699) with: {'learning_rate': 0.01, 'max_depth': 1, 'n_estimators
': 700}
-2.191846 (0.035863) with: {'learning_rate': 0.01, 'max_depth': 2, 'n_estimators
': 100}
-0.822023 (0.020079) with: {'learning_rate': 0.01, 'max_depth': 2, 'n_estimators
': 300}
-0.766157 (0.023321) with: {'learning_rate': 0.01, 'max_depth': 2, 'n_estimators
': 500}
-0.753672 (0.023543) with: {'learning_rate': 0.01, 'max_depth': 2, 'n_estimators
': 700}
-2.150979 (0.030036) with: {'learning_rate': 0.01, 'max_depth': 3, 'n_estimators
': 100}
-0.793687 (0.020444) with: {'learning_rate': 0.01, 'max_depth': 3, 'n_estimators
': 300}
-0.752255 (0.023239) with: {'learning_rate': 0.01, 'max_depth': 3, 'n_estimators
': 500}
-0.747185 (0.023241) with: {'learning_rate': 0.01, 'max_depth': 3, 'n_estimators
': 700}
-0.766614 (0.023431) with: {'learning_rate': 0.1, 'max_depth': 1, 'n_estimators'
: 100}
-0.747297 (0.022442) with: {'learning_rate': 0.1, 'max_depth': 1, 'n_estimators'
: 300}
-0.746802 (0.022036) with: {'learning_rate': 0.1, 'max_depth': 1, 'n_estimators'
: 500}
-0.746790 (0.021930) with: {'learning_rate': 0.1, 'max_depth': 1, 'n_estimators'
: 700}
-0.749163 (0.023078) with: {'learning_rate': 0.1, 'max_depth': 2, 'n_estimators'
: 100}
-0.745507 (0.022560) with: {'learning_rate': 0.1, 'max_depth': 2, 'n_estimators'
: 300}
-0.745194 (0.022640) with: {'learning_rate': 0.1, 'max_depth': 2, 'n_estimators'
: 500}
```

```
In [0]: xgb_bsl = xgb.XGBRegressor(max_depth=3, learning_rate = 0.1, n_estimators=500, nthread=-1)
xgb_bsl
```

```
Out[0]: XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
  colsample_bytree=1, gamma=0, learning_rate=0.1, max_delta_step=0,
  max_depth=3, min_child_weight=1, missing=None, n_estimators=500,
  n_jobs=1, nthread=-1, objective='reg:linear', random_state=0,
  reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
  silent=True, subsample=1)
```

```
In [0]: train_results, test_results = run_xgboost(xgb_bsl, x_train, y_train, x_test, y_test)

# store the results in models_evaluations dictionaries
models_evaluation_train['xgb_bsl'] = train_results
models_evaluation_test['xgb_bsl'] = test_results

xgb.plot_importance(xgb_bsl)
plt.show()
```

Training the model..

Done. Time taken : 0:02:20.248262

Done


Evaluating the model with TRAIN data...

Evaluating Test data

TEST DATA

RMSE : 1.0890322448240302

MAPE : 35.13968692492444



4.4.2 Suprise BaselineModel

```
In [0]: from surprise import BaselineOnly
```

Predicted_rating : (baseline prediction)

- http://surprise.readthedocs.io/en/stable/basic_algorithms.html#surprise.prediction_algorithms.baseline_only.BaselineOnly

$$\hat{r}_{ui} = b_{ui} = \mu + b_u + b_i$$

- μ : Average of all trainings in training data.
- b_u : User bias
- b_i : Item bias (movie biases)

Optimization function (Least Squares Problem)

- http://surprise.readthedocs.io/en/stable/prediction_algorithms.html#baselines-estimates-configuration

$$\sum_{r_{ui} \in R_{train}} (r_{ui} - (\mu + b_u + b_i))^2 + \lambda (b_u^2 + b_i^2) . \text{ [mimimize } b_u, b_i]$$

```
In [0]: # options are to specify.., how to compute those user and item biases
bsl_options = {'method': 'sgd',
               'reg':0.01,
               'learning_rate': 0.001,
               'n_epochs':120
              }
bsl_algo = BaselineOnly(bsl_options=bsl_options)
# run this algorithm.., It will return the train and test results..
bsl_train_results, bsl_test_results = run_surprise(bsl_algo, trainset, testset, verbose=True)

# Just store these error metrics in our models_evaluation datastructure
models_evaluation_train['bsl_algo'] = bsl_train_results
models_evaluation_test['bsl_algo'] = bsl_test_results

Training the model...
Estimating biases using sgd...
Done. time taken : 0:00:19.165750

Evaluating the model with train data..
time taken : 0:00:04.750962
-----
Train Data
-----
RMSE : 0.8982370573392073

MAPE : 27.429673745139915

adding train results in the dictionary..

Evaluating for test data...
time taken : 0:00:00.456780
-----
Test Data
-----
RMSE : 1.0865215481719563

MAPE : 34.9957270093008

storing the test results in test dictionary...

-----
Total time taken to run this algorithm : 0:00:24.373492
```

4.4.3 XGBoost with initial 13 features + Surprise Baseline predictor

Updating Train Data

```
In [0]: # add our baseline_predicted value as our feature..
reg_train['bslpr'] = models_evaluation_train['bsl_algo']['predictions']
reg_train.head(2)
```

Out[0]:

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	smr3	smr4	smr5	UAvg	MAvg
0	174683	10	3.586988	5.0	4.0	4.0	3.0	4.0	3.0	5.0	3.0	3.0	2.0	3.882353	3.636364
1	233949	10	3.586988	4.0	4.0	5.0	1.0	5.0	2.0	3.0	3.0	3.0	3.0	2.692308	3.636364

Updating Test Data

```
In [0]: # add that baseline predicted ratings with Surprise to the test data as well
reg_test_df['bslpr'] = models_evaluation_test['bsl_algo']['predictions']

reg_test_df.head(2)
```

Out[0]:

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	smr3
0	808635	71	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988
1	898730	71	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988

```
In [0]: # prepare train data
x_train = reg_train.drop(['user', 'movie', 'rating'], axis=1)
y_train = reg_train['rating']

# Prepare Test data
x_test = reg_test_df.drop(['user', 'movie', 'rating'], axis=1)
y_test = reg_test_df['rating']

# initialize Our first XGBoost model...
start = datetime.now()

# Initialize Our first XGBoost model
xgb = xgb.XGBRegressor(nthread=-1)

# Perform cross validation
gscv = GridSearchCV(xgb,
                    param_grid = parameters,
                    scoring="neg_mean_squared_error",
                    cv = TimeSeriesSplit(n_splits=5),
                    n_jobs = -1,
                    verbose = 1)
gscv_result = gscv.fit(x_train, y_train)

# Summarize results
print("Best: %f using %s" % (gscv_result.best_score_, gscv_result.best_params_))
print()
means = gscv_result.cv_results_['mean_test_score']
stds = gscv_result.cv_results_['std_test_score']
params = gscv_result.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))

print("\nTime Taken: ", datetime.now() - start)
```


Fitting 5 folds for each of 36 candidates, totalling 180 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.  
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed: 19.6min  
[Parallel(n_jobs=-1)]: Done 180 out of 180 | elapsed: 117.6min finished
```

Best: -0.744999 using {'learning_rate': 0.1, 'max_depth': 3, 'n_estimators': 500}

-8.940064 (0.136602) with: {'learning_rate': 0.001, 'max_depth': 1, 'n_estimators': 100}
-6.322584 (0.104569) with: {'learning_rate': 0.001, 'max_depth': 1, 'n_estimators': 300}
-4.560439 (0.079736) with: {'learning_rate': 0.001, 'max_depth': 1, 'n_estimators': 500}
-3.375169 (0.062092) with: {'learning_rate': 0.001, 'max_depth': 1, 'n_estimators': 700}
-8.910045 (0.119083) with: {'learning_rate': 0.001, 'max_depth': 2, 'n_estimators': 100}
-6.268585 (0.086633) with: {'learning_rate': 0.001, 'max_depth': 2, 'n_estimators': 300}
-4.494120 (0.066408) with: {'learning_rate': 0.001, 'max_depth': 2, 'n_estimators': 500}
-3.303398 (0.052270) with: {'learning_rate': 0.001, 'max_depth': 2, 'n_estimators': 700}
-8.900899 (0.118331) with: {'learning_rate': 0.001, 'max_depth': 3, 'n_estimators': 100}
-6.245090 (0.084058) with: {'learning_rate': 0.001, 'max_depth': 3, 'n_estimators': 300}
-4.463697 (0.061850) with: {'learning_rate': 0.001, 'max_depth': 3, 'n_estimators': 500}
-3.267050 (0.045774) with: {'learning_rate': 0.001, 'max_depth': 3, 'n_estimators': 700}
-2.271043 (0.045529) with: {'learning_rate': 0.01, 'max_depth': 1, 'n_estimators': 100}
-0.897886 (0.022963) with: {'learning_rate': 0.01, 'max_depth': 1, 'n_estimators': 300}
-0.819336 (0.023904) with: {'learning_rate': 0.01, 'max_depth': 1, 'n_estimators': 500}
-0.790015 (0.023699) with: {'learning_rate': 0.01, 'max_depth': 1, 'n_estimators': 700}
-2.191846 (0.035863) with: {'learning_rate': 0.01, 'max_depth': 2, 'n_estimators': 100}
-0.822023 (0.020079) with: {'learning_rate': 0.01, 'max_depth': 2, 'n_estimators': 300}
-0.766157 (0.023321) with: {'learning_rate': 0.01, 'max_depth': 2, 'n_estimators': 500}
-0.753672 (0.023543) with: {'learning_rate': 0.01, 'max_depth': 2, 'n_estimators': 700}
-2.150979 (0.030036) with: {'learning_rate': 0.01, 'max_depth': 3, 'n_estimators': 100}
-0.793687 (0.020444) with: {'learning_rate': 0.01, 'max_depth': 3, 'n_estimators': 300}
-0.752262 (0.023245) with: {'learning_rate': 0.01, 'max_depth': 3, 'n_estimators': 500}
-0.747189 (0.023244) with: {'learning_rate': 0.01, 'max_depth': 3, 'n_estimators': 700}
-0.766614 (0.023431) with: {'learning_rate': 0.1, 'max_depth': 1, 'n_estimators': 100}
-0.747315 (0.022474) with: {'learning_rate': 0.1, 'max_depth': 1, 'n_estimators': 300}
-0.746831 (0.022080) with: {'learning_rate': 0.1, 'max_depth': 1, 'n_estimators': 500}
-0.746798 (0.021941) with: {'learning_rate': 0.1, 'max_depth': 1, 'n_estimators': 700}
-0.749163 (0.023078) with: {'learning_rate': 0.1, 'max_depth': 2, 'n_estimators': 100}
-0.745505 (0.022538) with: {'learning_rate': 0.1, 'max_depth': 2, 'n_estimators': 300}
-0.745347 (0.022602) with: {'learning_rate': 0.1, 'max_depth': 2, 'n_estimators': 500}

```
In [0]: import xgboost as xgb
```

```
In [0]: xgb_bsl = xgb.XGBRegressor(max_depth=3, learning_rate = 0.1, n_estimators=500, nthread=-1)
xgb_bsl
```

```
Out[0]: XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                    colsample_bytree=1, gamma=0, learning_rate=0.1, max_delta_step=0,
                    max_depth=3, min_child_weight=1, missing=None, n_estimators=500,
                    n_jobs=1, nthread=-1, objective='reg:linear', random_state=0,
                    reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
                    silent=True, subsample=1)
```

```
In [0]: train_results, test_results = run_xgboost(xgb_bsl, x_train, y_train, x_test, y_test)

# store the results in models_evaluations dictionaries
models_evaluation_train['xgb_bsl'] = train_results
models_evaluation_test['xgb_bsl'] = test_results

xgb.plot_importance(xgb_bsl)
plt.show()
```

Training the model..

Done. Time taken : 0:03:09.679057

Done


Evaluating the model with TRAIN data...

Evaluating Test data

TEST DATA

RMSE : 1.0891181427027241

MAPE : 35.13135164276489



4.4.4 Surprise KNNBaseline predictor

```
In [0]: from surprise import KNNBaseline
```

- KNN BASELINE

- http://surprise.readthedocs.io/en/stable/knn_inspired.html#surprise.prediction_algorithms.knns.KNNBaseline
(http://surprise.readthedocs.io/en/stable/knn_inspired.html#surprise.prediction_algorithms.knns.KNNBaseline)

- PEARSON_BASELINE SIMILARITY

- http://surprise.readthedocs.io/en/stable/similarities.html#surprise.similarities.pearson_baseline
(http://surprise.readthedocs.io/en/stable/similarities.html#surprise.similarities.pearson_baseline)

- SHRINKAGE

- 2.2 Neighborhood Models in <http://courses.ischool.berkeley.edu/i290-dm/s11/SECURE/a1-koren.pdf>
(<http://courses.ischool.berkeley.edu/i290-dm/s11/SECURE/a1-koren.pdf>)

- predicted Rating : (**based on User-User similarity**)

$$\hat{r}_{ui} = b_{ui} + \frac{\sum_{v \in N_i^k(u)} \text{sim}(u, v) \cdot (r_{vi} - b_{vi})}{\sum_{v \in N_i^k(u)} \text{sim}(u, v)}$$

- b_{ui} - Baseline prediction of (user, movie) rating
- $N_i^k(u)$ - Set of **K similar** users (neighbours) of **user (u)** who rated **movie(i)**
- $\text{sim}(u, v)$ - **Similarity** between users **u** and **v**
 - Generally, it will be cosine similarity or Pearson correlation coefficient.
 - But we use **shrunk Pearson-baseline correlation coefficient**, which is based on the pearsonBaseline similarity (we take base line predictions instead of mean rating of user/item)

- Predicted rating (based on Item Item similarity):

$$\hat{r}_{ui} = b_{ui} + \frac{\sum_{j \in N_u^k(i)} \text{sim}(i, j) \cdot (r_{uj} - b_{uj})}{\sum_{j \in N_u^k(i)} \text{sim}(i, j)}$$

- **Notations follows same as above (user user based predicted rating)**

4.4.4.1 Surprise KNNBaseline with user user similarities

```
In [0]: # we specify , how to compute similarities and what to consider with sim_options to
our algorithm
sim_options = {'user_based' : True,
               'name': 'pearson_baseline',
               'shrinkage': 100,
               'min_support': 2
               }

# we keep other parameters like regularization parameter and learning_rate as default values.
bsl_options = {'method': 'sgd'}

knn_bsl_u = KNNBaseline(k=40, sim_options = sim_options, bsl_options = bsl_options)
knn_bsl_u_train_results, knn_bsl_u_test_results = run_surprise(knn_bsl_u, trainset,
testset, verbose=True)

# Just store these error metrics in our models_evaluation datastructure
models_evaluation_train['knn_bsl_u'] = knn_bsl_u_train_results
models_evaluation_test['knn_bsl_u'] = knn_bsl_u_test_results
```

```
Training the model...
Estimating biases using sgd...
Computing the pearson_baseline similarity matrix...
Done computing similarity matrix.
Done. time taken : 3:21:13.767364

Evaluating the model with train data..
time taken : 0:15:14.838640
-----
Train Data
-----
RMSE : 0.4495623286931499

MAPE : 12.69889589980268

adding train results in the dictionary..

Evaluating for test data...
time taken : 0:00:00.798002
-----
Test Data
-----
RMSE : 1.0865005562678032

MAPE : 35.02325234274119

storing the test results in test dictionary...

-----
Total time taken to run this algorithm : 3:36:29.513762
```

4.4.4.2 Surprise KNNBaseline with movie movie similarities

```

In [0]: # we specify , how to compute similarities and what to consider with sim_options to
our algorithm

# 'user_based' : Fals => this considers the similarities of movies instead of users

sim_options = {'user_based' : False,
               'name': 'pearson_baseline',
               'shrinkage': 100,
               'min_support': 2
               }

# we keep other parameters like regularization parameter and learning_rate as default values.
bsl_options = {'method': 'sgd'}

knn_bsl_m = KNNBaseline(k=40, sim_options = sim_options, bsl_options = bsl_options)

knn_bsl_m_train_results, knn_bsl_m_test_results = run_surprise(knn_bsl_m, trainset,
testset, verbose=True)

# Just store these error metrics in our models_evaluation datastructure
models_evaluation_train['knn_bsl_m'] = knn_bsl_m_train_results
models_evaluation_test['knn_bsl_m'] = knn_bsl_m_test_results

```

Training the model...

Estimating biases using sgd...

Computing the pearson_baseline similarity matrix...

Done computing similarity matrix.

Done. time taken : 0:00:14.860019

Evaluating the model with train data..

time taken : 0:01:33.572011

Train Data

RMSE : 0.49544620093528796

MAPE : 13.87789147087551

adding train results in the dictionary..

Evaluating for test data...

time taken : 0:00:00.624896

Test Data

RMSE : 1.0868914468761874

MAPE : 35.02725521759712

storing the test results in test dictionary...

Total time taken to run this algorithm : 0:01:49.056926

4.4.5 XGBoost with initial 13 features + Surprise Baseline predictor + KNNBaseline predictor

- First we will run XGBoost with predictions from both KNN's (that uses User_User and Item_Item similarities along with our previous features.
- Then we will run XGBoost with just predictions from both knn models and predictions from our baseline model.

Preparing Train data

```
In [0]: # add the predicted values from both knns to this dataframe
reg_train['knn_bsl_u'] = models_evaluation_train['knn_bsl_u']['predictions']
reg_train['knn_bsl_m'] = models_evaluation_train['knn_bsl_m']['predictions']

reg_train.head(2)
```

Out[0]:

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	smr3	smr4	smr5	UAvg	MAvg
0	174683	10	3.586988	5.0	4.0	4.0	3.0	4.0	3.0	5.0	3.0	3.0	2.0	3.882353	3.636364
1	233949	10	3.586988	4.0	4.0	5.0	1.0	5.0	2.0	3.0	3.0	3.0	3.0	2.692308	3.636364

Preparing Test data

```
In [0]: reg_test_df['knn_bsl_u'] = models_evaluation_test['knn_bsl_u']['predictions']
reg_test_df['knn_bsl_m'] = models_evaluation_test['knn_bsl_m']['predictions']

reg_test_df.head(2)
```

Out[0]:

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	smr3
0	808635	71	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988
1	898730	71	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988


```
In [0]: # prepare the train data....
x_train = reg_train.drop(['user', 'movie', 'rating'], axis=1)
y_train = reg_train['rating']

# prepare the train data....
x_test = reg_test_df.drop(['user', 'movie', 'rating'], axis=1)
y_test = reg_test_df['rating']

start = datetime.now()

# Initialize Our first XGBoost model
model = xgb.XGBRegressor(nthread=-1)

# Perform cross validation
gscv = GridSearchCV(model,
                    param_grid = parameters,
                    scoring="neg_mean_squared_error",
                    cv = TimeSeriesSplit(n_splits=5),
                    n_jobs = -1,
                    verbose = 1)
gscv_result = gscv.fit(x_train, y_train)

# Summarize results
print("Best: %f using %s" % (gscv_result.best_score_, gscv_result.best_params_))
print()
means = gscv_result.cv_results_['mean_test_score']
stds = gscv_result.cv_results_['std_test_score']
params = gscv_result.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))
```

Fitting 5 folds for each of 36 candidates, totalling 180 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.  
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed: 22.1min  
[Parallel(n_jobs=-1)]: Done 180 out of 180 | elapsed: 120.5min finished
```

Best: -0.745209 using {'learning_rate': 0.1, 'max_depth': 3, 'n_estimators': 300}

-8.940064 (0.136602) with: {'learning_rate': 0.001, 'max_depth': 1, 'n_estimators': 100}
-6.322584 (0.104569) with: {'learning_rate': 0.001, 'max_depth': 1, 'n_estimators': 300}
-4.560439 (0.079736) with: {'learning_rate': 0.001, 'max_depth': 1, 'n_estimators': 500}
-3.375169 (0.062092) with: {'learning_rate': 0.001, 'max_depth': 1, 'n_estimators': 700}
-8.910045 (0.119083) with: {'learning_rate': 0.001, 'max_depth': 2, 'n_estimators': 100}
-6.268585 (0.086633) with: {'learning_rate': 0.001, 'max_depth': 2, 'n_estimators': 300}
-4.494120 (0.066408) with: {'learning_rate': 0.001, 'max_depth': 2, 'n_estimators': 500}
-3.303398 (0.052270) with: {'learning_rate': 0.001, 'max_depth': 2, 'n_estimators': 700}
-8.900899 (0.118331) with: {'learning_rate': 0.001, 'max_depth': 3, 'n_estimators': 100}
-6.245090 (0.084058) with: {'learning_rate': 0.001, 'max_depth': 3, 'n_estimators': 300}
-4.463697 (0.061850) with: {'learning_rate': 0.001, 'max_depth': 3, 'n_estimators': 500}
-3.267050 (0.045774) with: {'learning_rate': 0.001, 'max_depth': 3, 'n_estimators': 700}
-2.271043 (0.045529) with: {'learning_rate': 0.01, 'max_depth': 1, 'n_estimators': 100}
-0.897886 (0.022963) with: {'learning_rate': 0.01, 'max_depth': 1, 'n_estimators': 300}
-0.819336 (0.023904) with: {'learning_rate': 0.01, 'max_depth': 1, 'n_estimators': 500}
-0.790015 (0.023699) with: {'learning_rate': 0.01, 'max_depth': 1, 'n_estimators': 700}
-2.191846 (0.035863) with: {'learning_rate': 0.01, 'max_depth': 2, 'n_estimators': 100}
-0.822023 (0.020079) with: {'learning_rate': 0.01, 'max_depth': 2, 'n_estimators': 300}
-0.766157 (0.023321) with: {'learning_rate': 0.01, 'max_depth': 2, 'n_estimators': 500}
-0.753672 (0.023543) with: {'learning_rate': 0.01, 'max_depth': 2, 'n_estimators': 700}
-2.150979 (0.030036) with: {'learning_rate': 0.01, 'max_depth': 3, 'n_estimators': 100}
-0.793687 (0.020444) with: {'learning_rate': 0.01, 'max_depth': 3, 'n_estimators': 300}
-0.752257 (0.023243) with: {'learning_rate': 0.01, 'max_depth': 3, 'n_estimators': 500}
-0.747195 (0.023238) with: {'learning_rate': 0.01, 'max_depth': 3, 'n_estimators': 700}
-0.766614 (0.023431) with: {'learning_rate': 0.1, 'max_depth': 1, 'n_estimators': 100}
-0.747312 (0.022469) with: {'learning_rate': 0.1, 'max_depth': 1, 'n_estimators': 300}
-0.746835 (0.022091) with: {'learning_rate': 0.1, 'max_depth': 1, 'n_estimators': 500}
-0.746810 (0.021962) with: {'learning_rate': 0.1, 'max_depth': 1, 'n_estimators': 700}
-0.749180 (0.023109) with: {'learning_rate': 0.1, 'max_depth': 2, 'n_estimators': 100}
-0.745436 (0.022539) with: {'learning_rate': 0.1, 'max_depth': 2, 'n_estimators': 300}
-0.745413 (0.022615) with: {'learning_rate': 0.1, 'max_depth': 2, 'n_estimators': 500}

```
In [0]: xgb_knn_bsl = xgb.XGBRegressor(max_depth=3, learning_rate = 0.1, n_estimators=300, nthread=-1)
xgb_knn_bsl
```

```
Out[0]: XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
  colsample_bytree=1, gamma=0, learning_rate=0.1, max_delta_step=0,
  max_depth=3, min_child_weight=1, missing=None, n_estimators=300,
  n_jobs=1, nthread=-1, objective='reg:linear', random_state=0,
  reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
  silent=True, subsample=1)
```

```
In [0]: train_results, test_results = run_xgboost(xgb_knn_bsl, x_train, y_train, x_test, y_test)

# store the results in models_evaluations dictionaries
models_evaluation_train['xgb_knn_bsl'] = train_results
models_evaluation_test['xgb_knn_bsl'] = test_results

xgb.plot_importance(xgb_knn_bsl)
plt.show()
```

Training the model..

Done. Time taken : 0:01:35.173342

Done


Evaluating the model with TRAIN data...

Evaluating Test data

TEST DATA

RMSE : 1.088749005744821

MAPE : 35.188974153659295



4.4.6 Matrix Factorization Techniques

4.4.6.1 SVD Matrix Factorization User Movie interactions

```
In [0]: from surprise import SVD
```

http://surprise.readthedocs.io/en/stable/matrix_factorization.html#surprise.prediction_algorithms.matrix_factorization.SVD
[\(http://surprise.readthedocs.io/en/stable/matrix_factorization.html#surprise.prediction_algorithms.matrix_factorization.SVD\)](http://surprise.readthedocs.io/en/stable/matrix_factorization.html#surprise.prediction_algorithms.matrix_factorization.SVD)

- Predicted Rating :

$$- \hat{r}_{ui} = \mu + b_u + b_i + q_i^T p_u$$

- q_i - Representation of item(movie) in latent factor space

- p_u - Representation of user in new latent factor space

- A BASIC MATRIX FACTORIZATION MODEL in [https://datajobs.com/data-science-repo/Recommender-Systems-\[Netflix\].pdf](https://datajobs.com/data-science-repo/Recommender-Systems-[Netflix].pdf) ([https://datajobs.com/data-science-repo/Recommender-Systems-\[Netflix\].pdf](https://datajobs.com/data-science-repo/Recommender-Systems-[Netflix].pdf))

- Optimization problem with user item interactions and regularization (to avoid overfitting)

$$- \sum_{r_{ui} \in R_{train}} \left(r_{ui} - \hat{r}_{ui} \right)^2 +$$

$$\lambda (b_i^2 + b_u^2 + \|q_i\|^2 + \|p_u\|^2)$$

```
In [0]: # initialize the model
svd = SVD(n_factors=100, biased=True, random_state=15, verbose=True)
svd_train_results, svd_test_results = run_surprise(svd, trainset, testset, verbose=True)

# Just store these error metrics in our models_evaluation datastructure
models_evaluation_train['svd'] = svd_train_results
models_evaluation_test['svd'] = svd_test_results
```

Training the model...

Processing epoch 0

Processing epoch 1

Processing epoch 2

Processing epoch 3

Processing epoch 4

Processing epoch 5

Processing epoch 6

Processing epoch 7

Processing epoch 8

Processing epoch 9

Processing epoch 10

Processing epoch 11

Processing epoch 12

Processing epoch 13

Processing epoch 14

Processing epoch 15

Processing epoch 16

Processing epoch 17

Processing epoch 18

Processing epoch 19

Done. time taken : 0:00:39.893902

Evaluating the model with train data..

time taken : 0:00:06.312768

Train Data

RMSE : 0.6702496300850848

MAPE : 19.93649200841313

adding train results in the dictionary..

Evaluating for test data...

time taken : 0:00:00.671748

Test Data

RMSE : 1.0860031195730506

MAPE : 34.94819349312387

storing the test results in test dictionary...

Total time taken to run this algorithm : 0:00:46.878418

4.4.6.2 SVD Matrix Factorization with implicit feedback from user (user rated movies)

```
In [0]: from surprise import SVDpp
```

- ----> 2.5 Implicit Feedback in <http://courses.ischool.berkeley.edu/i290-dm/s11/SECURE/a1-koren.pdf>
(<http://courses.ischool.berkeley.edu/i290-dm/s11/SECURE/a1-koren.pdf>)

- Predicted Rating :

$$\hat{r}_{ui} = \mu + b_u + b_i + q_i^T \left(p_u + \frac{1}{|I_u|} \sum_{j \in I_u} y_j \right)$$

- I_u --- the set of all items rated by user u
- y_j --- Our new set of item factors that capture implicit ratings.

- Optimization problem with user item interactions and regularization (to avoid overfitting)

$$\min \sum_{\{ui\} \in R_{\text{train}}} (r_{ui} - \hat{r}_{ui})^2 + \lambda (b_i^2 + b_u^2 + \|q_i\|^2 + \|p_u\|^2 + \|y_j\|^2)$$


```
In [0]: # initialize the model
svdpp = SVDpp(n_factors=50, random_state=15, verbose=True)
svdpp_train_results, svdpp_test_results = run_surprise(svdpp, trainset, testset, verbose=True)

# Just store these error metrics in our models_evaluation datastructure
models_evaluation_train['svdpp'] = svdpp_train_results
models_evaluation_test['svdpp'] = svdpp_test_results
```

Training the model...

processing epoch 0
processing epoch 1
processing epoch 2
processing epoch 3
processing epoch 4
processing epoch 5
processing epoch 6
processing epoch 7
processing epoch 8
processing epoch 9
processing epoch 10
processing epoch 11
processing epoch 12
processing epoch 13
processing epoch 14
processing epoch 15
processing epoch 16
processing epoch 17
processing epoch 18
processing epoch 19

Done. time taken : 0:25:19.547805

Evaluating the model with train data..

time taken : 0:01:02.080985

Train Data

RMSE : 0.6581255901775523

MAPE : 19.083570120018518

adding train results in the dictionary..

Evaluating for test data...

time taken : 0:00:00.484273

Test Data

RMSE : 1.0862780572420558

MAPE : 34.909882014758175

storing the test results in test dictionary...

Total time taken to run this algorithm : 0:26:22.113063

4.4.7 XgBoost with 13 features + Surprise Baseline + Surprise KNNbaseline + MF Techniques

Preparing Train data

```
In [0]: # add the predicted values from both knns to this dataframe
reg_train['svd'] = models_evaluation_train['svd']['predictions']
reg_train['svdpp'] = models_evaluation_train['svdpp']['predictions']

reg_train.head(2)
```

Out [0]:

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	smr3	smr4	smr5	UAvg	MAvg
0	174683	10	3.586988	5.0	4.0	4.0	3.0	4.0	3.0	5.0	3.0	3.0	2.0	3.882353	3.636364
1	233949	10	3.586988	4.0	4.0	5.0	1.0	5.0	2.0	3.0	3.0	3.0	3.0	2.692308	3.636364

Preparing Test data

```
In [0]: reg_test_df['svd'] = models_evaluation_test['svd']['predictions']
reg_test_df['svdpp'] = models_evaluation_test['svdpp']['predictions']

reg_test_df.head(2)
```

Out [0]:

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	smr3
0	808635	71	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988
1	898730	71	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988	3.586988

```
In [0]: # prepare x_train and y_train
x_train = reg_train.drop(['user', 'movie', 'rating'], axis=1)
y_train = reg_train['rating']

# prepare test data
x_test = reg_test_df.drop(['user', 'movie', 'rating'], axis=1)
y_test = reg_test_df['rating']

start = datetime.now()

# Initialize Our first XGBoost model
model = xgb.XGBRegressor(nthread=-1)

# Perform cross validation
gscv = GridSearchCV(model,
                    param_grid = parameters,
                    scoring="neg_mean_squared_error",
                    cv = TimeSeriesSplit(n_splits=5),
                    n_jobs = -1,
                    verbose = 1)
gscv_result = gscv.fit(x_train, y_train)

# Summarize results
print("Best: %f using %s" % (gscv_result.best_score_, gscv_result.best_params_))
print()
means = gscv_result.cv_results_['mean_test_score']
stds = gscv_result.cv_results_['std_test_score']
params = gscv_result.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))

print("\nTime Taken: ", datetime.now() - start)
```

Fitting 5 folds for each of 36 candidates, totalling 180 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.  
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed: 28.4min  
[Parallel(n_jobs=-1)]: Done 180 out of 180 | elapsed: 168.0min finished
```

Best: -0.745372 using {'learning_rate': 0.1, 'max_depth': 3, 'n_estimators': 300}

-8.940064 (0.136602) with: {'learning_rate': 0.001, 'max_depth': 1, 'n_estimators': 100}
-6.322584 (0.104569) with: {'learning_rate': 0.001, 'max_depth': 1, 'n_estimators': 300}
-4.560439 (0.079736) with: {'learning_rate': 0.001, 'max_depth': 1, 'n_estimators': 500}
-3.375169 (0.062092) with: {'learning_rate': 0.001, 'max_depth': 1, 'n_estimators': 700}
-8.910045 (0.119083) with: {'learning_rate': 0.001, 'max_depth': 2, 'n_estimators': 100}
-6.268585 (0.086633) with: {'learning_rate': 0.001, 'max_depth': 2, 'n_estimators': 300}
-4.494120 (0.066408) with: {'learning_rate': 0.001, 'max_depth': 2, 'n_estimators': 500}
-3.303398 (0.052270) with: {'learning_rate': 0.001, 'max_depth': 2, 'n_estimators': 700}
-8.900899 (0.118331) with: {'learning_rate': 0.001, 'max_depth': 3, 'n_estimators': 100}
-6.245090 (0.084058) with: {'learning_rate': 0.001, 'max_depth': 3, 'n_estimators': 300}
-4.463697 (0.061850) with: {'learning_rate': 0.001, 'max_depth': 3, 'n_estimators': 500}
-3.267050 (0.045774) with: {'learning_rate': 0.001, 'max_depth': 3, 'n_estimators': 700}
-2.271043 (0.045529) with: {'learning_rate': 0.01, 'max_depth': 1, 'n_estimators': 100}
-0.897886 (0.022963) with: {'learning_rate': 0.01, 'max_depth': 1, 'n_estimators': 300}
-0.819336 (0.023904) with: {'learning_rate': 0.01, 'max_depth': 1, 'n_estimators': 500}
-0.790015 (0.023699) with: {'learning_rate': 0.01, 'max_depth': 1, 'n_estimators': 700}
-2.191846 (0.035863) with: {'learning_rate': 0.01, 'max_depth': 2, 'n_estimators': 100}
-0.822023 (0.020079) with: {'learning_rate': 0.01, 'max_depth': 2, 'n_estimators': 300}
-0.766157 (0.023321) with: {'learning_rate': 0.01, 'max_depth': 2, 'n_estimators': 500}
-0.753672 (0.023543) with: {'learning_rate': 0.01, 'max_depth': 2, 'n_estimators': 700}
-2.150979 (0.030036) with: {'learning_rate': 0.01, 'max_depth': 3, 'n_estimators': 100}
-0.793687 (0.020444) with: {'learning_rate': 0.01, 'max_depth': 3, 'n_estimators': 300}
-0.752259 (0.023237) with: {'learning_rate': 0.01, 'max_depth': 3, 'n_estimators': 500}
-0.747194 (0.023236) with: {'learning_rate': 0.01, 'max_depth': 3, 'n_estimators': 700}
-0.766614 (0.023431) with: {'learning_rate': 0.1, 'max_depth': 1, 'n_estimators': 100}
-0.747312 (0.022470) with: {'learning_rate': 0.1, 'max_depth': 1, 'n_estimators': 300}
-0.746840 (0.022108) with: {'learning_rate': 0.1, 'max_depth': 1, 'n_estimators': 500}
-0.746814 (0.021988) with: {'learning_rate': 0.1, 'max_depth': 1, 'n_estimators': 700}
-0.749180 (0.023109) with: {'learning_rate': 0.1, 'max_depth': 2, 'n_estimators': 100}
-0.745497 (0.022593) with: {'learning_rate': 0.1, 'max_depth': 2, 'n_estimators': 300}
-0.745480 (0.022647) with: {'learning_rate': 0.1, 'max_depth': 2, 'n_estimators': 500}

```
In [0]: xgb_final = xgb.XGBRegressor(max_depth=3, learning_rate = 0.1, n_estimators=300, nthread=-1)
xgb_final

Out[0]: XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
  colsample_bytree=1, gamma=0, learning_rate=0.1, max_delta_step=0,
  max_depth=3, min_child_weight=1, missing=None, n_estimators=300,
  n_jobs=1, nthread=-1, objective='reg:linear', random_state=0,
  reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
  silent=True, subsample=1)
```

```
In [0]: train_results, test_results = run_xgboost(xgb_final, x_train, y_train, x_test, y_test)

# store the results in models_evaluations dictionaries
models_evaluation_train['xgb_final'] = train_results
models_evaluation_test['xgb_final'] = test_results

xgb.plot_importance(xgb_final)
plt.show()
```

Training the model..

Done. Time taken : 0:01:44.972457

Done


Evaluating the model with TRAIN data...

Evaluating Test data

TEST DATA

RMSE : 1.0891599523508655

MAPE : 35.12646240961147



In [0]:

4.4.8 XgBoost with Surprise Baseline + Surprise KNNbaseline + MF Techniques

```
In [0]: # prepare train data
x_train = reg_train[['knn_bsl_u', 'knn_bsl_m', 'svd', 'svdpp']]
y_train = reg_train['rating']

# test data
x_test = reg_test_df[['knn_bsl_u', 'knn_bsl_m', 'svd', 'svdpp']]
y_test = reg_test_df['rating']

start = datetime.now()

# Initialize Our first XGBoost model
model = xgb.XGBRegressor(nthread=-1)

# Perform cross validation
gscv = GridSearchCV(model,
                    param_grid = parameters,
                    scoring="neg_mean_squared_error",
                    cv = TimeSeriesSplit(n_splits=5),
                    n_jobs = -1,
                    verbose = 1)
gscv_result = gscv.fit(x_train, y_train)

# Summarize results
print("Best: %f using %s" % (gscv_result.best_score_, gscv_result.best_params_))
print()
means = gscv_result.cv_results_['mean_test_score']
stds = gscv_result.cv_results_['std_test_score']
params = gscv_result.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))

print("\nTime Taken: ",datetime.now() - start)
```


Fitting 5 folds for each of 36 candidates, totalling 180 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.  
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed: 13.8min  
[Parallel(n_jobs=-1)]: Done 180 out of 180 | elapsed: 89.8min finished
```

Best: -1.171865 using {'learning_rate': 0.01, 'max_depth': 1, 'n_estimators': 700}

-8.963209 (0.136970) with: {'learning_rate': 0.001, 'max_depth': 1, 'n_estimators': 100}
-6.392774 (0.116929) with: {'learning_rate': 0.001, 'max_depth': 1, 'n_estimators': 300}
-4.670259 (0.099834) with: {'learning_rate': 0.001, 'max_depth': 1, 'n_estimators': 500}
-3.515961 (0.085541) with: {'learning_rate': 0.001, 'max_depth': 1, 'n_estimators': 700}
-8.963206 (0.136959) with: {'learning_rate': 0.001, 'max_depth': 2, 'n_estimators': 100}
-6.392728 (0.116905) with: {'learning_rate': 0.001, 'max_depth': 2, 'n_estimators': 300}
-4.670230 (0.099787) with: {'learning_rate': 0.001, 'max_depth': 2, 'n_estimators': 500}
-3.515905 (0.085520) with: {'learning_rate': 0.001, 'max_depth': 2, 'n_estimators': 700}
-8.963209 (0.136942) with: {'learning_rate': 0.001, 'max_depth': 3, 'n_estimators': 100}
-6.392755 (0.116862) with: {'learning_rate': 0.001, 'max_depth': 3, 'n_estimators': 300}
-4.670291 (0.099755) with: {'learning_rate': 0.001, 'max_depth': 3, 'n_estimators': 500}
-3.515990 (0.085470) with: {'learning_rate': 0.001, 'max_depth': 3, 'n_estimators': 700}
-2.445828 (0.068581) with: {'learning_rate': 0.01, 'max_depth': 1, 'n_estimators': 100}
-1.194459 (0.035786) with: {'learning_rate': 0.01, 'max_depth': 1, 'n_estimators': 300}
-1.172220 (0.034604) with: {'learning_rate': 0.01, 'max_depth': 1, 'n_estimators': 500}
-1.171865 (0.034555) with: {'learning_rate': 0.01, 'max_depth': 1, 'n_estimators': 700}
-2.445742 (0.068594) with: {'learning_rate': 0.01, 'max_depth': 2, 'n_estimators': 100}
-1.194504 (0.035828) with: {'learning_rate': 0.01, 'max_depth': 2, 'n_estimators': 300}
-1.172317 (0.034653) with: {'learning_rate': 0.01, 'max_depth': 2, 'n_estimators': 500}
-1.171985 (0.034612) with: {'learning_rate': 0.01, 'max_depth': 2, 'n_estimators': 700}
-2.445784 (0.068580) with: {'learning_rate': 0.01, 'max_depth': 3, 'n_estimators': 100}
-1.194518 (0.035891) with: {'learning_rate': 0.01, 'max_depth': 3, 'n_estimators': 300}
-1.172375 (0.034704) with: {'learning_rate': 0.01, 'max_depth': 3, 'n_estimators': 500}
-1.172076 (0.034672) with: {'learning_rate': 0.01, 'max_depth': 3, 'n_estimators': 700}
-1.171891 (0.034556) with: {'learning_rate': 0.1, 'max_depth': 1, 'n_estimators': 100}
-1.171990 (0.034594) with: {'learning_rate': 0.1, 'max_depth': 1, 'n_estimators': 300}
-1.172046 (0.034612) with: {'learning_rate': 0.1, 'max_depth': 1, 'n_estimators': 500}
-1.172092 (0.034620) with: {'learning_rate': 0.1, 'max_depth': 1, 'n_estimators': 700}
-1.172106 (0.034647) with: {'learning_rate': 0.1, 'max_depth': 2, 'n_estimators': 100}
-1.172518 (0.034816) with: {'learning_rate': 0.1, 'max_depth': 2, 'n_estimators': 300}
-1.172958 (0.034952) with: {'learning_rate': 0.1, 'max_depth': 2, 'n_estimators': 500}

```
In [0]: xgb_all_models = xgb.XGBRegressor(max_depth=1, learning_rate = 0.01, n_estimators=700,
      , nthread=-1)
      xgb_all_models
```

```
Out[0]: XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
      colsample_bytree=1, gamma=0, learning_rate=0.01, max_delta_step=0,
      max_depth=1, min_child_weight=1, missing=None, n_estimators=700,
      n_jobs=1, nthread=-1, objective='reg:linear', random_state=0,
      reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
      silent=True, subsample=1)
```

```
In [0]: train_results, test_results = run_xgboost(xgb_all_models, x_train, y_train, x_test,
      y_test)
```

```
# store the results in models_evaluations dictionaries
models_evaluation_train['xgb_all_models'] = train_results
models_evaluation_test['xgb_all_models'] = test_results
```

```
xgb.plot_importance(xgb_all_models)
plt.show()
```

Training the model..

Done. Time taken : 0:01:06.591248

Done


Evaluating the model with TRAIN data...

Evaluating Test data

TEST DATA

RMSE : 1.095123189648495

MAPE : 35.54329712868095



4.5 Comparision between all models

```
In [0]: # Saving our TEST_RESULTS into a dataframe so that you don't have to run it again
pd.DataFrame(models_evaluation_test).to_csv('sample/small/small_sample_results.csv'
)
models = pd.read_csv('sample/small/small_sample_results.csv', index_col=0)
models.loc['rmse'].sort_values()
```

```
Out[0]: svd                1.0860031195730506
svdpp              1.0862780572420558
knn_bsl_u          1.0865005562678032
knn_bsl_m          1.0868914468761874
xgb_knn_bsl        1.088749005744821
xgb_final          1.0891599523508655
xgb_all_models     1.095123189648495
Name: rmse, dtype: object
```

```
In [0]: from prettytable import PrettyTable
numbering = [1,2,3,4,5,6,7,8,9,10]
featurization = ['svd','knn_bsl_u','bsl_algo','knn_bsl_m','svdpp','xgb_final','xgb_bsl',
'first_algo','xgb_knn_bsl','xgb_all_models']
rmse=['1.0860031195730506','1.0865005562678032','1.0868914468761874','1.0865215481
719563',
'1.0868914468761874','1.0862780572420558','1.0891599523508655',
'1.0890322448240302','1.088749005744821','1.095123189648495' ]
ptable = PrettyTable()

# Adding columns
ptable.add_column("S.NO.",numbering)
ptable.add_column("MODEL",featurization)
ptable.add_column("RMSE",rmse)

# Printing the Table
print(ptable)
```

S.NO.	MODEL	RMSE
1	svd	1.0860031195730506
2	knn_bsl_u	1.0865005562678032
3	bsl_algo	1.0868914468761874
4	knn_bsl_m	1.0865215481719563
5	svdpp	1.0868914468761874
6	xgb_final	1.0862780572420558
7	xgb_bsl	1.0891599523508655
8	first_algo	1.0890322448240302
9	xgb_knn_bsl	1.088749005744821
10	xgb_all_models	1.095123189648495

Procedure

- 1.Merging movies with users and their rating in single dataframe.
- 2.Sorting the dataframe by date and removing any duplicate values.
- 3.performing Exploratory Data Analysis on it,so that we will able to visualise distribution of the ratings,avg rating of the movie or avg rating given by the users to the movie.
- 4.After that we we split our data in train and test which is in ratio of 80:20and try to to EDA on it.
- 5.Creating sparse matrix from train and test data and analysing the sparsity.
- 6.Compute User-User Similarity matrix and Movie-Movie similarity matrix.
- 7.Getting sample sparse matrices from train and test data.
- 8.Feurizing the data by selecting the top 5 ratings given by similar users to a particular user and also top 5 ratings given to similar movies with respect to particular movie.
- 9.work with different machine learning models and compare results.

Conclusion

- 1.Due to high computational cost, I have completed this case study on (18000,3000) training dataset and (9000,1500) testing dataset.
- 2.Every regressor model is hyper tuned for optimal parameters.
- 3.SVD model showed good result among all the models we tried.